

Performance evaluation, capacity planning and reliability

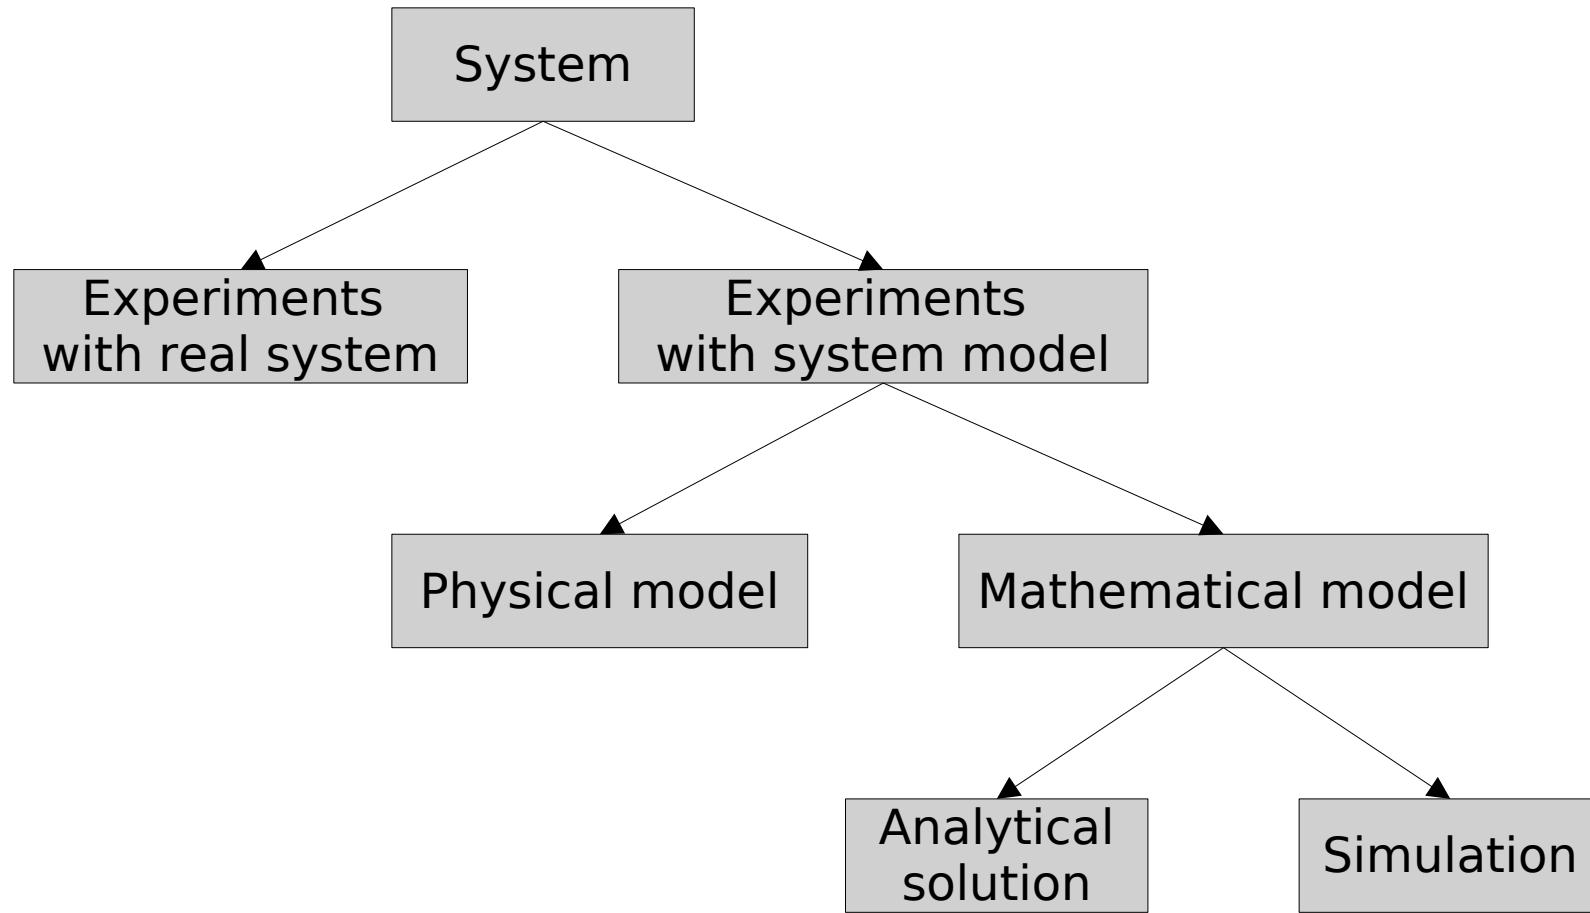
UNIMORE

UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA



Performance evaluation

Possible approaches



Possible approaches

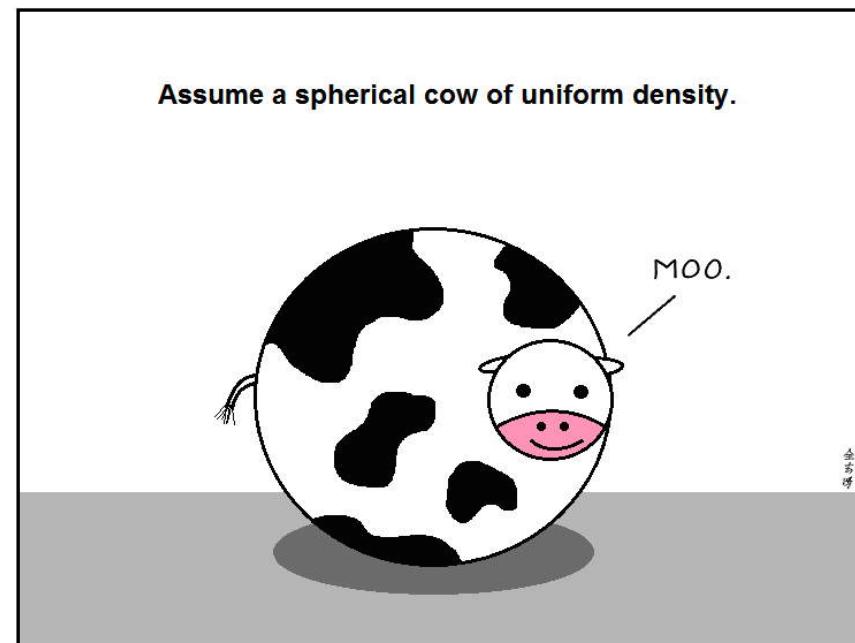
- Real system
 - Full-scale system
 - Testing and Benchmarking
- Physical model
 - Small scale prototype
 - Testing and benchmarking
 - Scaling up results
 - May be tricky in case of non-linear scalability

Possible approaches

- **Mathematical model**
 - Equations that describe system behavior
 - Easy to handle
 - What-if scenarios, comparison fo alternatives
 - Part of management algorithms
 - Simplified system description
- **Simulator**
 - Software that implement system behavior
 - More complex and realistic than mathematical model

Comparison of approaches

- **Actual system vs. System model**
 - Easier to do/Less dangerous (if system exists)
 - Less expensive than building a real system
- **Physical vs. Mathematical model**
 - Easier and less expensive to implement
 - Easier to modify
 - Mathematical model must be **accurate enough**
 - Remember the Spherical Cow
- **Analytical solution vs. Simulation**
 - Analytical is more **precise**
 - Simulation is more **flexible**



UNIMORE

UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA

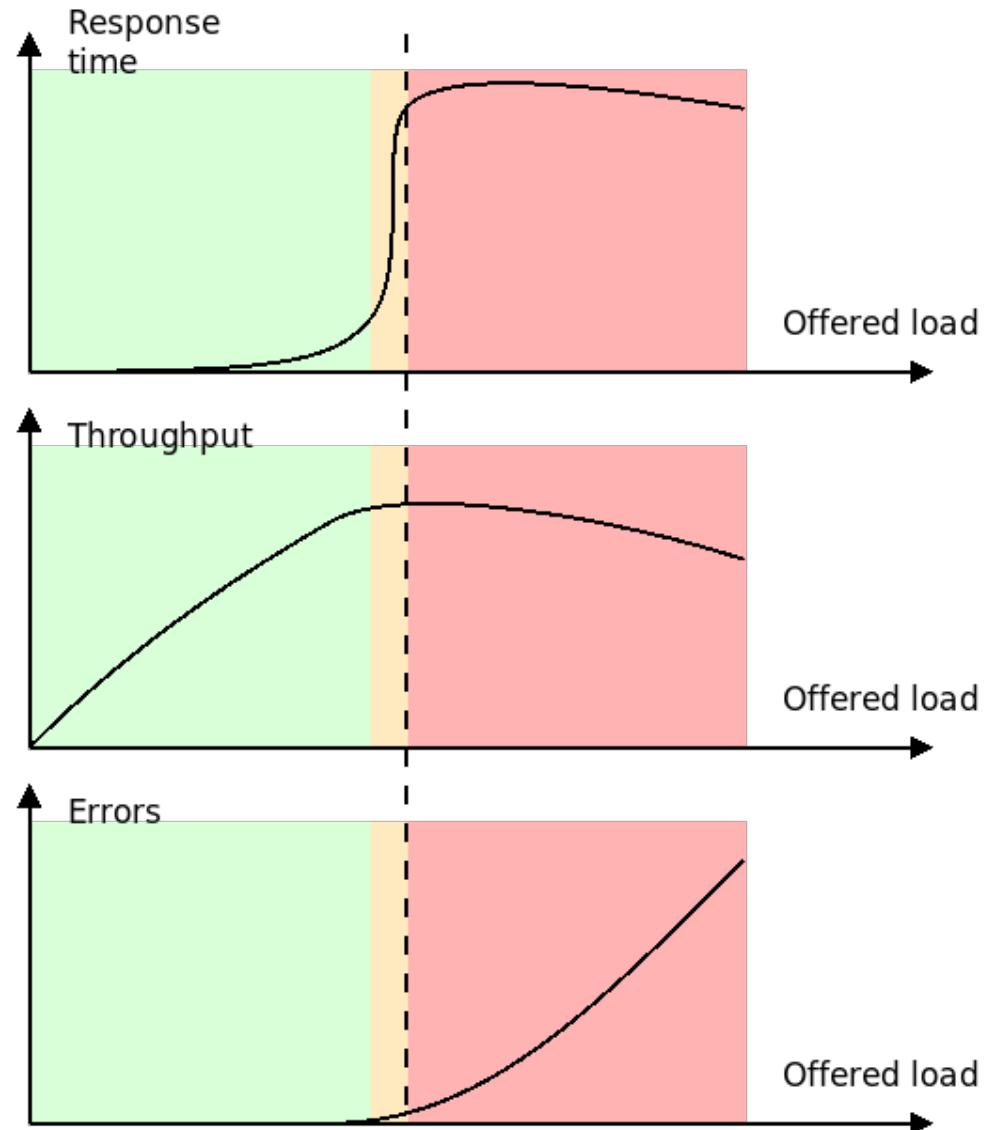


Motivational examples

Performance parameters

- **Response time:**
 - Elapsed time between
 - Beginning of request
 - End of Response
- **Throughput:**
 - Number of requests satisfied within a time unit
- **Error rate:**
 - Fraction of requests not served correctly
- Parameters depend on current load

Performance parameters



Operating conditions

- **OK/Underutilized** (Green)
 - Response time is low
 - Response time scarcely correlated with load
 - Throughput increases linearly with offered load
 - Error rate remains low
- **Knee region** (Yellow)
 - Response time increases dramatically with load
 - Throughput saturates
 - Error rate starts to grow

Operating conditions

- Overload/Thrashing (Red)
 - Response time explodes
 - Response time limited by timeout
 - Throughput is constant (overload)
 - Throughput drops (thrashing)
 - High error rate
- Our goal:
 - remaining at the end of the green zone

The cost of errors

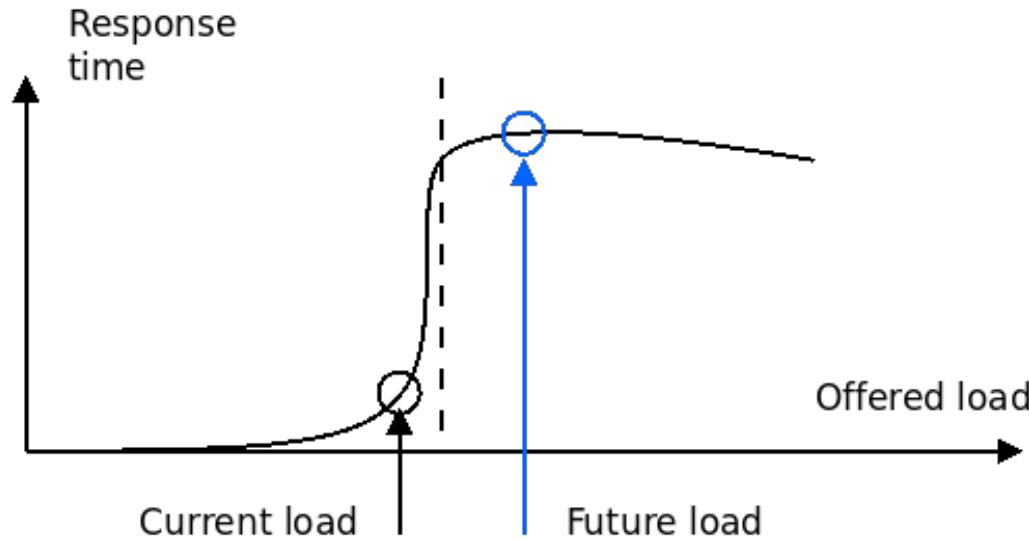
- Configuration errors have a cost
- A mis-configured system can be:
 - **Poor performing**
 - 60% slowdown observed in cloud deployments
 - A single bottleneck can slow down a huge chain of services
 - **Expensive**
 - Performance are OK, but cost is huge
 - 10× cost increase observed in cloud deployments

Capacity planning

- A system may be running fine
 - Today
 - **Tomorrow?**
- What about workload evolution?
- What if the load increases by $X\%$?
 - Are we still providing acceptable performance
 - If not, how do we cope with this?

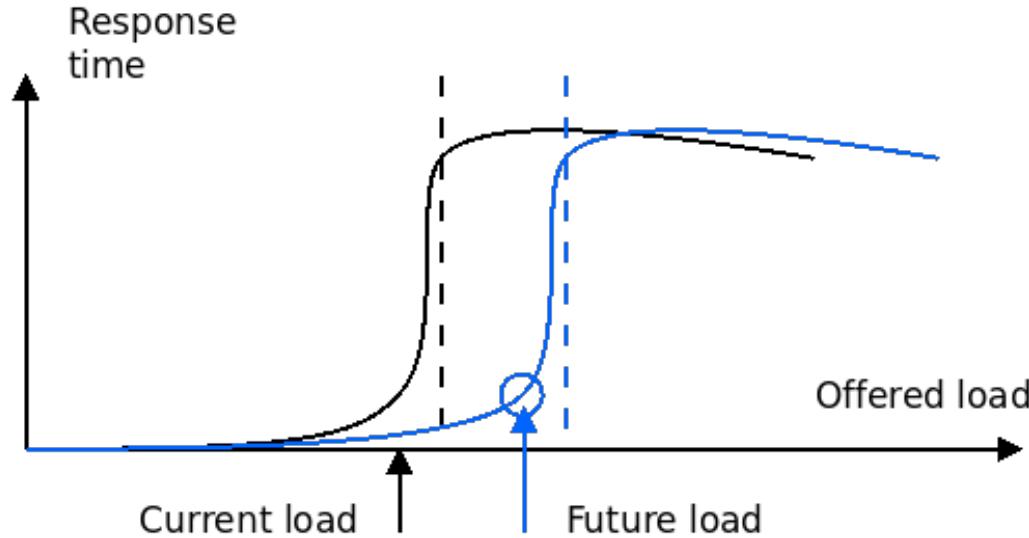


Load evolves



- **Current situation:**
 - Ideal operating condition
- **Future situation**
 - High risk of overload

Capacity planning



- Scaling is required
- New provisioning
 - How much more processing power do we need?
 - How to design the new infrastructure

Elasticity done well

- Cloud computing is elastic
 - Can purchase new computing power as needed
- We still need to **understand performance**:
 - How many VMs?
 - Which size of VMs?
 - For how much time?
 - What resource is the performance bottleneck?
- Cloud without understanding performance ~~can be is~~
 - Tricky
 - **Expensive**
- **Would risk your credit card?**



Another (unfortunate) case

- Scaling based on a model
 - Distributed infrastructure
 - Compute number of replicas to meet **SLA**
- Use of a **(over-)simplified** model
 - **M/M/1** queuing model

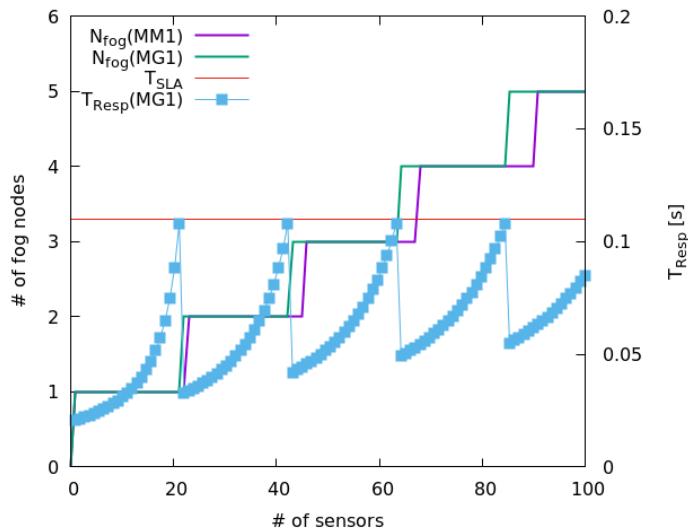
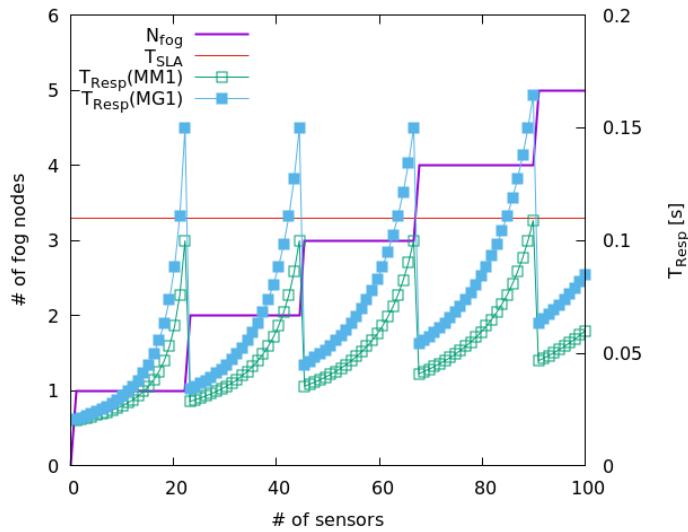
$$N = \sum_{j \in \mathcal{F}} E_j \geq \left\lceil \frac{\Lambda}{\bar{\mu}} \cdot \frac{K - 1}{K} \right\rceil$$

- Neglecting variance of service time variable
- What happens if the service time is highly variant
 - Common scenario: **high variance**
 - Should use log-normal distribution to describe service

Another (unfortunate) case

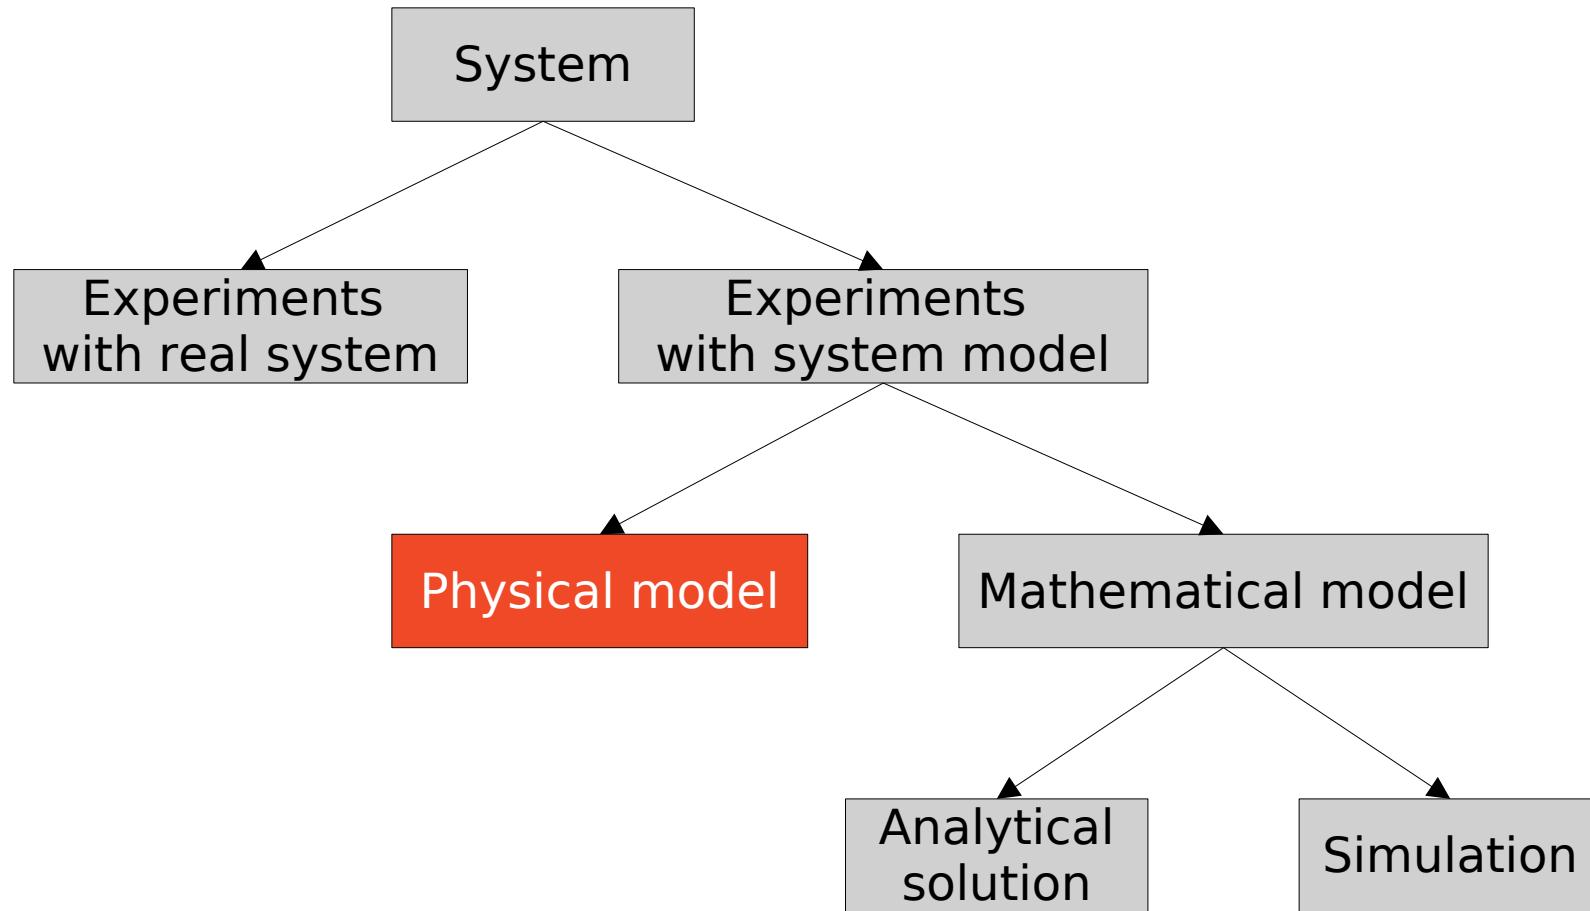
- Wrong model leads to wrong results
 - Expected (green) → OK
 - Real (blue)
→ **SLA violations**
- Should use correct model
 - Based on Pollaczek-Khinchin formula
 - Correct scaling

$$N \geq \left\lceil \frac{\Lambda}{\bar{\mu}} \cdot \frac{\text{CoV}^2 - 2K - 1}{2K - 2} \right\rceil$$



Real system testing

Possible approaches



Types of test

- **Functional** testing
 - System behaves as expected
- **Activity** testing
 - System works under normal conditions
- **Endurance** testing
 - System can scale under (limited) pressure
- **Stress** testing
 - Limits of the system
- **(Benchmarking)**
 - Performance under standardized load

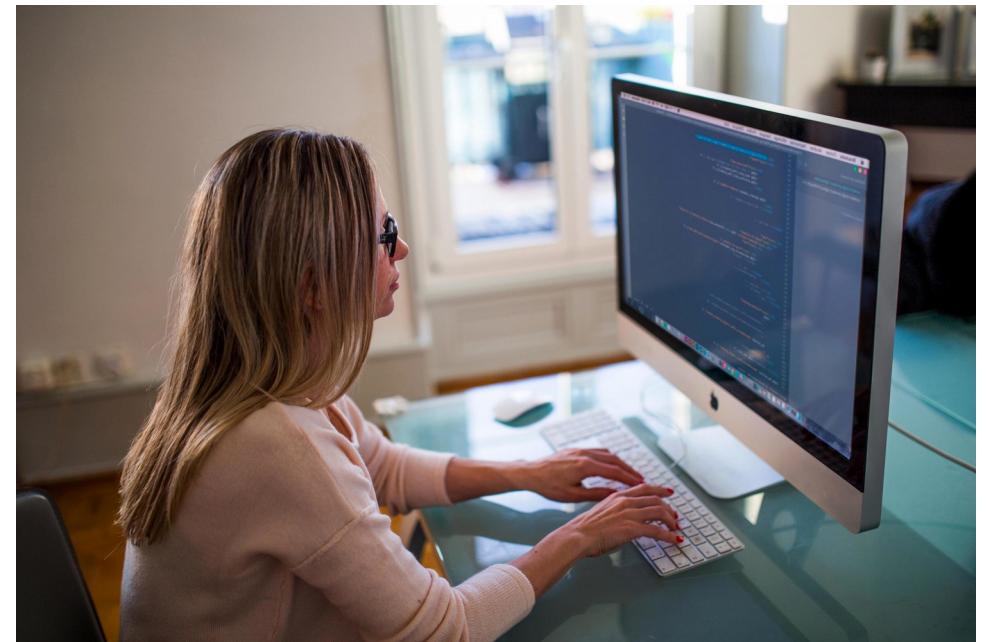
Functional testing

- Functional testing
- Definition of expected behavior
 - Test suite with known results
- Execution of tests
 - Checking expected results
 - Checking error catching
 - Checking for timeouts
 - Bugs/Regressions
- Example:
 - Unit testing



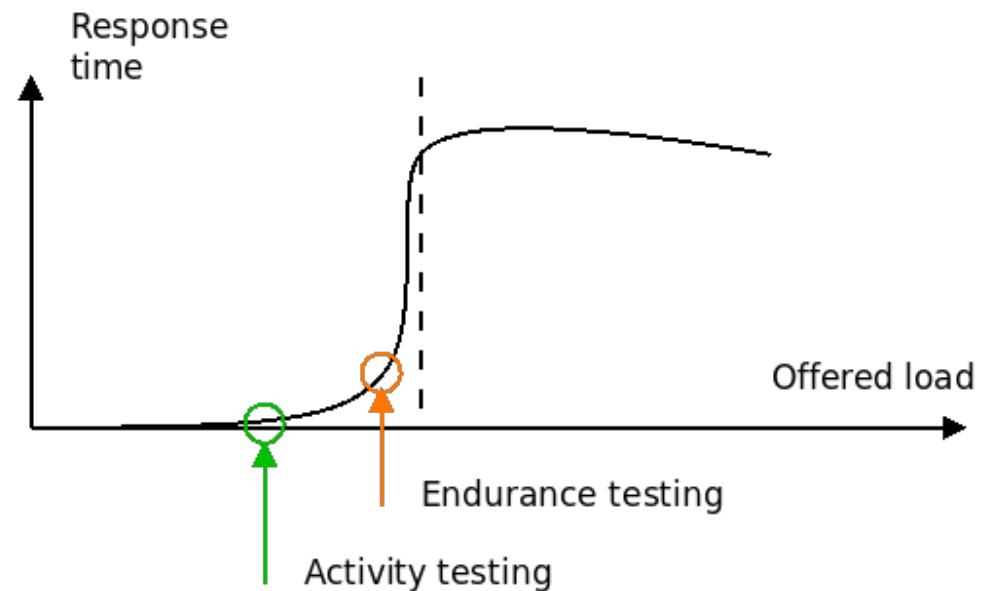
Activity testing

- Activity testing
- Typical workload
- Normal operation set
- Identification of problems
 - Beyond unit testing
- Example
 - Testing from users
 - Small production deployment



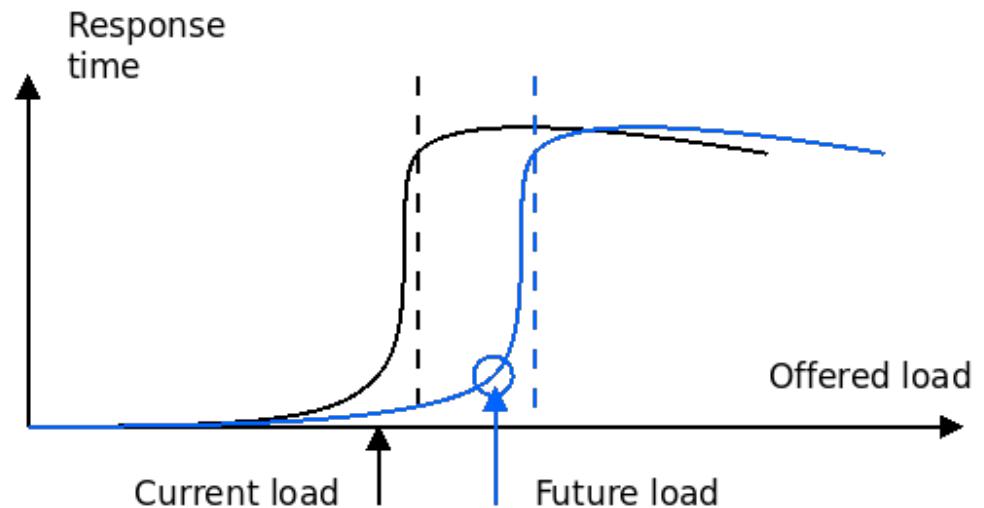
Endurance testing

- Endurance testing
- High load → system under stress
- Main focus
 - Is the system still working on rush hours?
 - What-If scenarios
 - *what happens if load increases by 20%?*



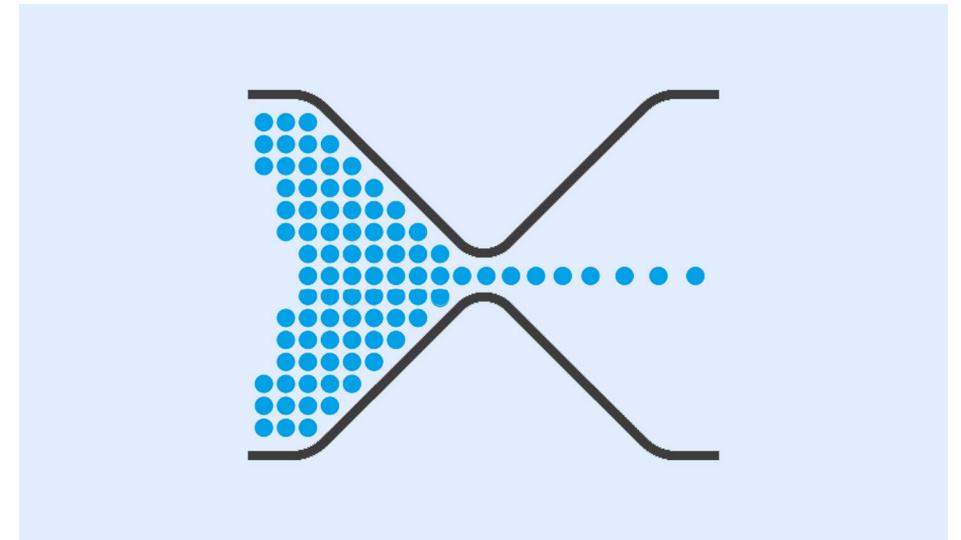
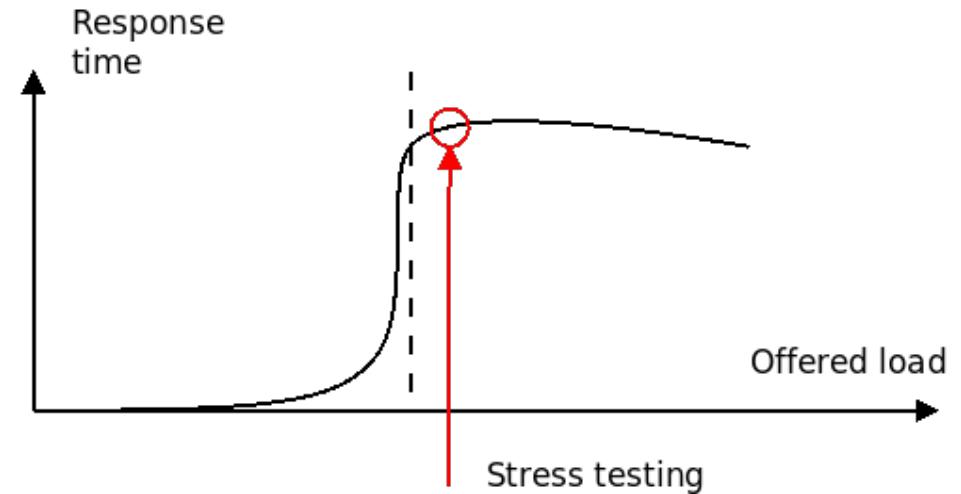
Endurance testing

- Endurance testing
- Goals
 - Evaluate scalability
 - Input for capacity planning
- System is under stress
 - *But it is expected to cope with the new conditions*



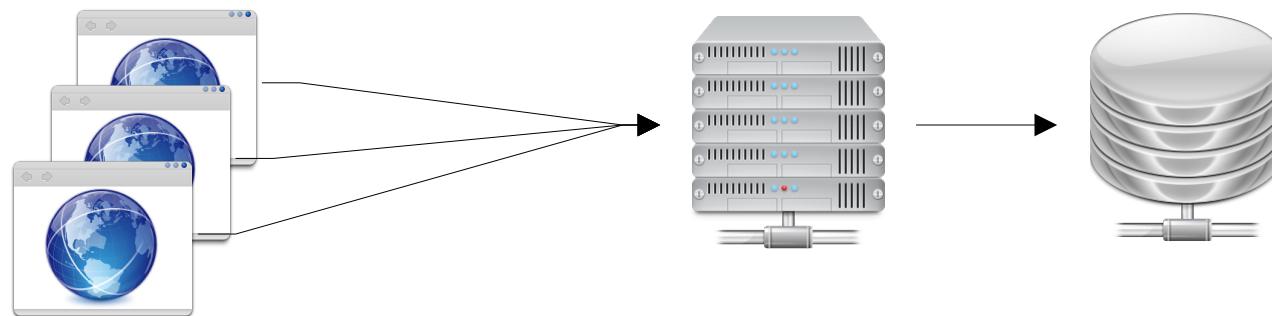
Stress testing

- Stress testing
- System is taken to the limit
 - We aim to create failures
- Main focus
 - What is the maximum load?
 - Where are the system bottlenecks?
 - Presence of trashing?



Example

- Example:
- 2-tier Web server
 - DBMS: up to 100 req/s
 - Web server: up to 500 req/s
 - 50% of Web page requests require access to DB
- What is the **maximum throughput**?
- Where is the **bottleneck**?

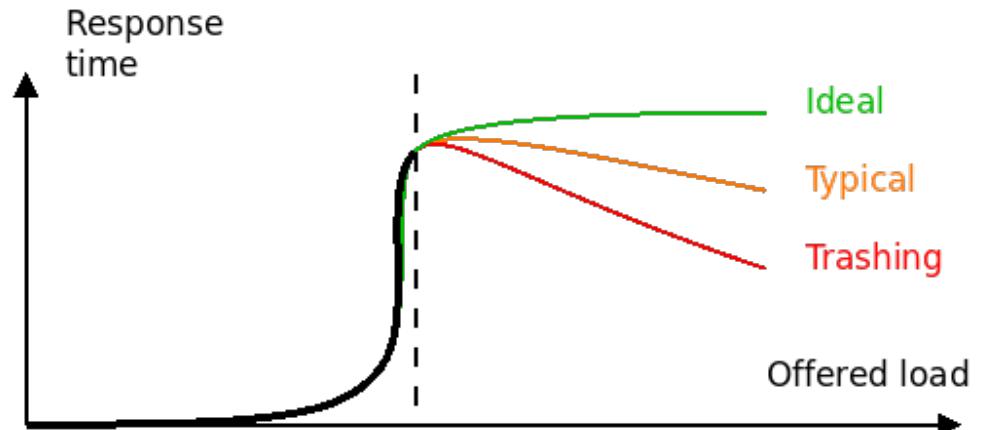


Example

- **Example:**
- Incoming load of 500 req/s
 - Web server saturates
 - DBMS receives $500 * 0.5 = 250$ req/s > 100 req/s
- Incoming load of 200 req/s
 - Web server utilization: 60%
 - DBMS is saturated (load = 100 req/s)
- Conclusion
 - Max load is 200 req/s
 - Bottleneck is DBMS

Trashing

- Stress testing
- Possible behavior
 - Ideal saturation
 - Trashing
 - Intermediate behavior
- Impact of thrashing
 - Cache locality disrupted
 - Resource contention



Benchmarking

- Performance evaluation requires a **workload**
- How can I **compare** my working condition with other scenarios?
- How can I **select the best option** of every experiment under different workload conditions?
- **Need for a standard**

Benchmarking

- Dictionary definition
 - *A point of reference by which something can be measured*
 - *(not necessarily related to computers)*
- Computer industry definition
 - *A set of conditions against which a product or system is measured*
 - *A set of performance criteria which a product is expected to meet*
- Critical concepts:
 - Point of reference
 - Benchmark is related to measurement

Benchmarking

- Load testing
- Evaluating performance
- Realistic workload conditions
- Refers to specific application scenario
- Results cannot be compared outside the specific application
- Benchmarking
- Evaluating performance
- Using workload conditions
 - Well-defined
 - Standardized
- Results can be compared between different systems

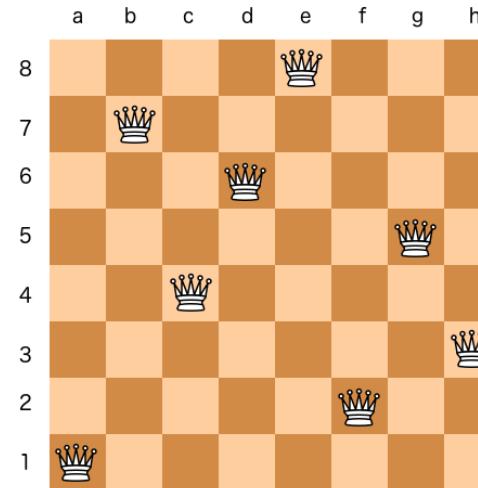
Benchmarking

- Types of benchmark
 - Toy benchmark
 - Real programs
 - Synthetic benchmarks
 - Benchmark based on real applications
- Not very representative
- Slightly better, prone to abuse
- Much better

Toy benchmark

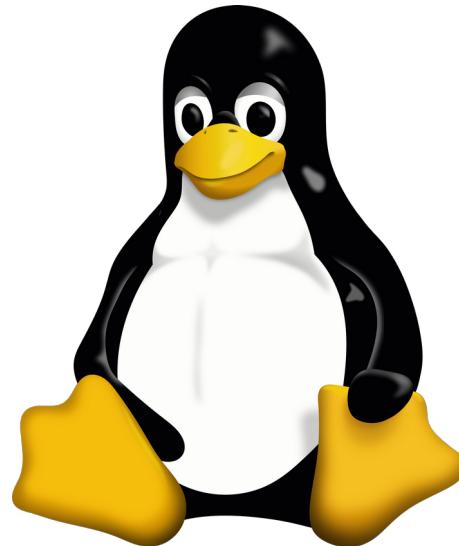
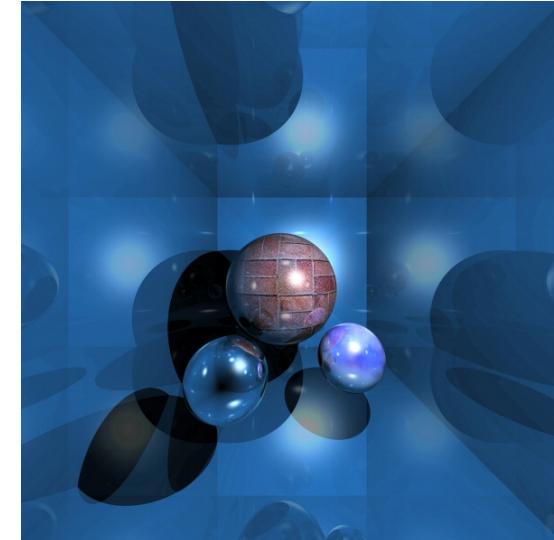
- Toy benchmark
- Small algorithms
- Examples:
 - Quicksort
 - Matrix operations
 - 8 queens problem
- Simple to build
- Easy to run
- Not very representative
 - *Who runs 8-queen problem in daily work?*

$$U = \begin{bmatrix} u_{1,1} & u_{1,2} & u_{1,3} & \dots & u_{1,n} \\ u_{2,1} & u_{2,2} & u_{2,3} & \dots & u_{2,n} \\ \vdots & \ddots & \ddots & \ddots & \vdots \\ 0 & & & \ddots & u_{n-1,n} \\ & & & & u_{n,n} \end{bmatrix}$$



Real programs

- Real programs
- Run a well known software on a platform
- Time for the task is a measure of system performance
- Examples
 - Ray tracing SW
 - Compiling the linux kernel

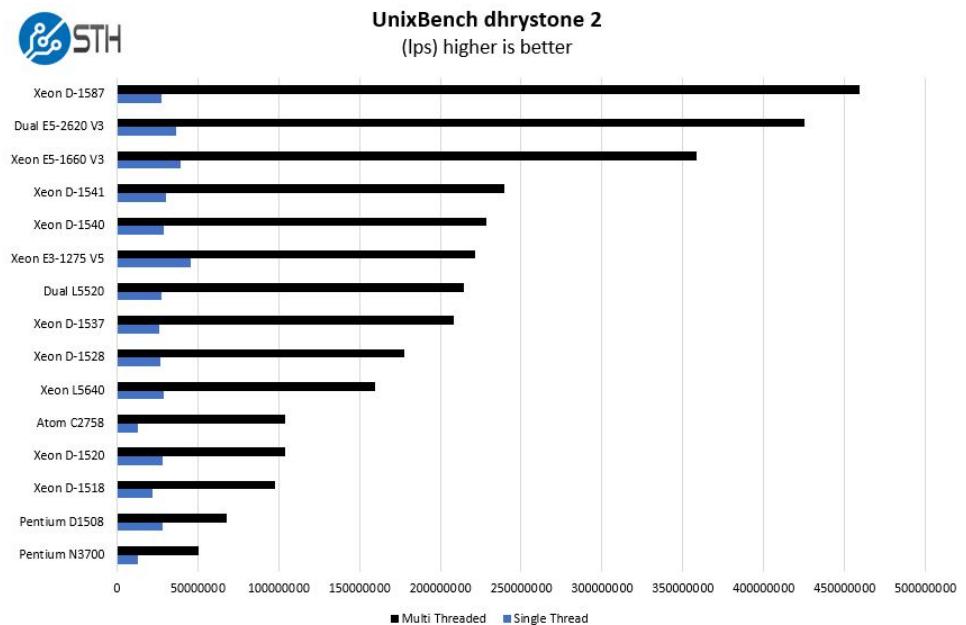


Real programs

- Some common **problems**
- Not every software works on every platform
 - **Porting** issues of gcc toolchain
- Optimization problems
 - **Driver issues** for some GPU
- Not clear if the **workload** is
 - Significant
 - General to draw meaningful conclusions
- If multiple SW are used
 - How to **weight** them properly?

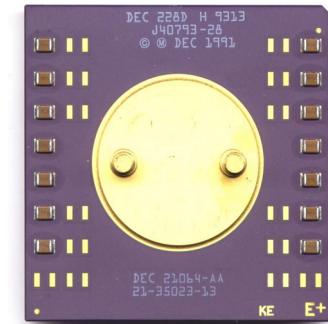
Synthetic benchmark

- Synthetic benchmarks
- Based on common operations in software
- Typically focused on some limited aspect of HW architecture
- Dhrystone
 - Integer math performance
- Whetstone
 - Floating point performance



Benchmark abuse

- Cases of **benchmark abuse**
 - Marketing impact of benchmark results
- **Hardware** architecture **optimized** for benchmark performance
 - Multiply-and-sum instruction (Power PC architecture)
- **Compiler optimization**
 - Unroll-loops in benchmark code
 - Jump optimization (Alpha CPU)
- Benchmark can be
 - Binary (pre-compiled)
 - Source (to compile)



Benchmark abuse

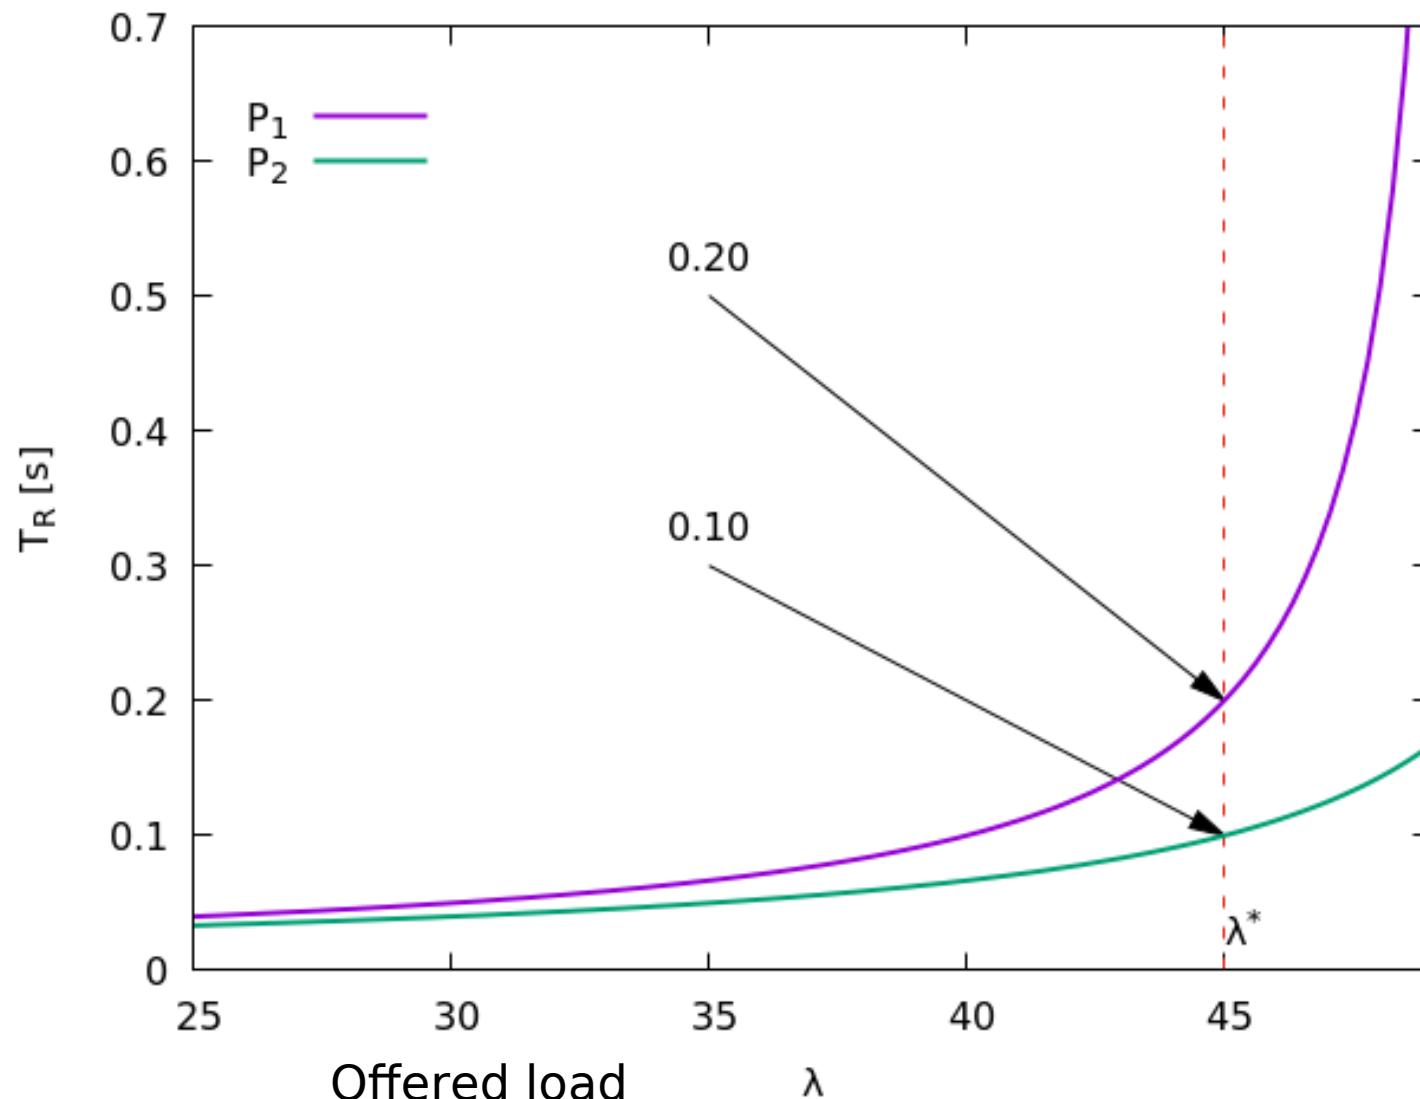
- Some common tricks
- Use **home-brewed** benchmark, give no reference
- Show only a **subset of suites** with better performance
- **Modify benchmark** source code
- Use **only part** of the benchmark tests
- Use **timing** of measures to manipulate **overheads**
 - E.g., discard first runs to mask cache misses
 - Measure only some overheads that make competitors look bad
- **Everything should be clearly stated and motivated**

A dirty trick

- Let's turn Mr. Hyde for an example
- We want to sell a new CPU
 - Slightly faster than competitors
 - Expensive
- Need good advertisement
 - Pay 50% more for a 10% performance increase
- We can *tune* our benchmark
 - When system is heavily loaded small changes in power make the difference in response time
 - Pay 50% more for 2× faster



A dirty trick



A dirty trick

- We will see the model detail later
- Lessons to bring home
 - Understanding performance model is important
 - Knowing benchmark test conditions is vital
- Now back to Dr. Jekyll mode



Real applications

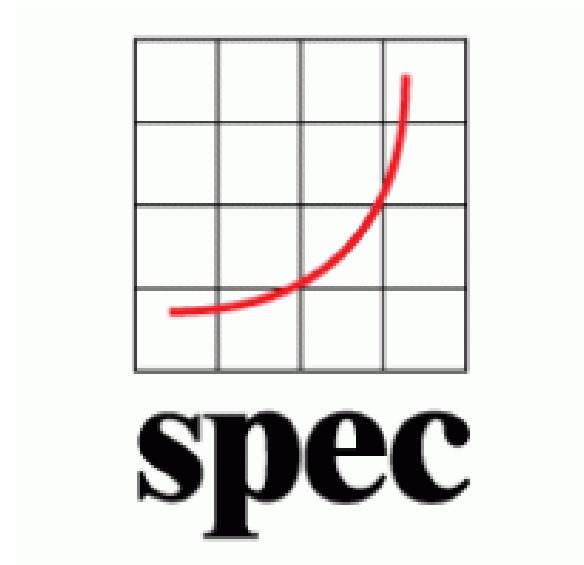
- Benchmark based on **real applications**
- Goals:
 - Make benchmark representative of **real problems**
 - Provide **open source** code and clear performance measurements
- Compare with **previous types of benchmark**
 - Full application (not just algorithms)
 - Representative problems (not just kernel compiling)
 - More than ALU testing (not just Int or FP math)
 - Fully standardized scenarios

Standards

- Old-school benchmarks
 - SPEC
 - TPC
- More cloud-friendly benchmarks
 - SPEC
 - Yahoo YCSB
 - PerfKit Benchmark

SPEC Web

- Developed by SPEC
 - System Performance Evaluation Consortium
 - Non-profit group developing several benchmark
- Representative workload
 - Web requests based on log traces
 - Logs from NCSA, HP, ...



- Workload
 - 4 classes of files
 - Poisson file size distribution in classes
- Scaling
 - N. directory in data set ×2 as data set size ×4
 - Requests evenly distributed on dirs
- Emulated Users model
 - On-off model
 - Think time for pages

Fraction [%]	Size
35%	0 - 1 KB
50%	1 - 10 KB
14%	10 - 100 KB
1%	100 KB - 1 MB

- Support for levels of burstiness
- Architecture:
 - One controller
 - Multiple clients
 - One server

- Evolution of benchmark
- Different directory size
- Support for HTTP/1.1
 - Keep-alive requests (70%)
 - Cookies
- Support for dynamic content
- Support for advertisement based on user-id cookie
- Zipf file popularity model

Fraction [%]	Type
70.00%	Static GET
12.45%	Dynamic GET
12.60%	Dynamic GET+ Ad
4.80%	Dynamic POST
0.15%	Dynamic GET + CGI

- TPC:
 - Transaction Processing Performance Council
- Several classes of benchmarks:
- TPC-W:
 - Web-based benchmark
- TPC-D
 - Database benchmark
- Each benchmark is based on a test application



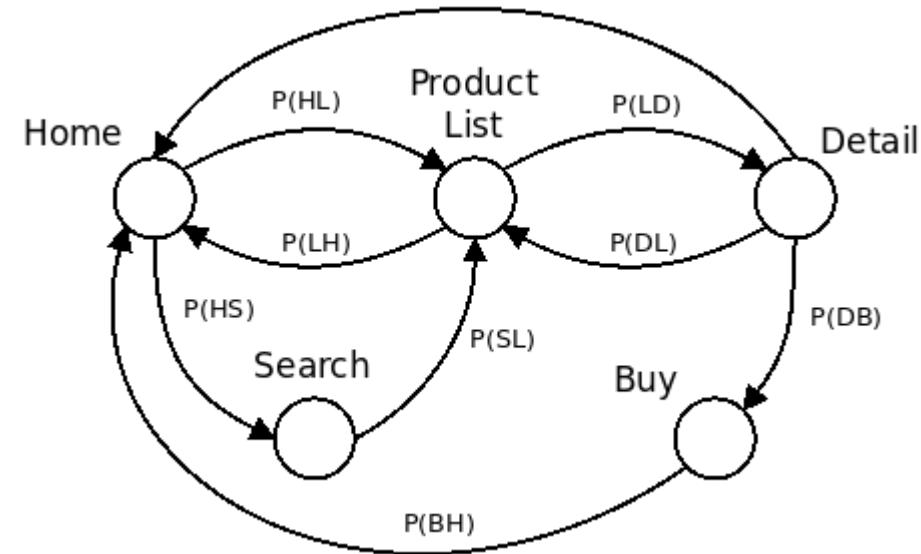
- **TPC-W**
- Dynamic Web site (changes with version)
 - Book store (like Amazon)
 - Auction system (like ebay)



- Auction system inspired by RUBiS benchmark
 - Rice University Bidding System
 - Auction in Java and PHP

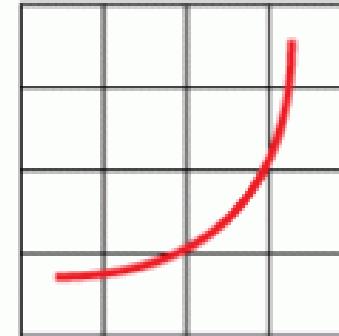


- Definition of workload
 - Several **dynamic web pages**
 - **Database size** (product catalog)
- Definition of **user behavior**
 - User defined as **stochastic FSM**
 - Each page is a state
 - Probability of passing from a page to another
- Population of users
- Multiple workload mixes (browsing/shopping/...)



SPEC-Cloud

- SPEC-Cloud
- IaaS-oriented benchmark
- Focus on
 - Scalability
 - Elasticity
- Test:
 - Scalability → load increases without creating errors
 - Elasticity → load changes, SLA preserved



spec

- YCSB
 - Yahoo Cloud Serving Benchmark
- Developed by Yahoo!
- Focus on Cassandra Database
- Available on github
- Test of transactions on a DB
 - Replicated
 - NoSQL



tjake/YCSB-test

Yahoo! Cloud Serving Benchmark



0 Contributors 0 Issues 2 Stars 1 Fork

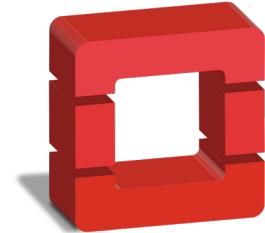
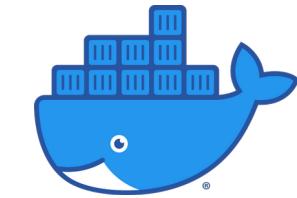


- Access to **Key-Value** storage
- Multi-phase benchmark
- **Baseline** phase
 - Run once the application
- **Elasticity+Scalability** phase
 - Re-run application with varying workload intensity
 - Collect measures
 - Compare with baseline

- Scalability metric
 - Throughput speedup compared to baseline
- Elasticity metric
 - Performance deviation from baseline
- Other metrics
 - **Instance provision time** (time to deploy new VM)
 - **Instance run success** (percentage of successful runs)
 - **Provisioning success** (percentage of new instance successfully deployed)

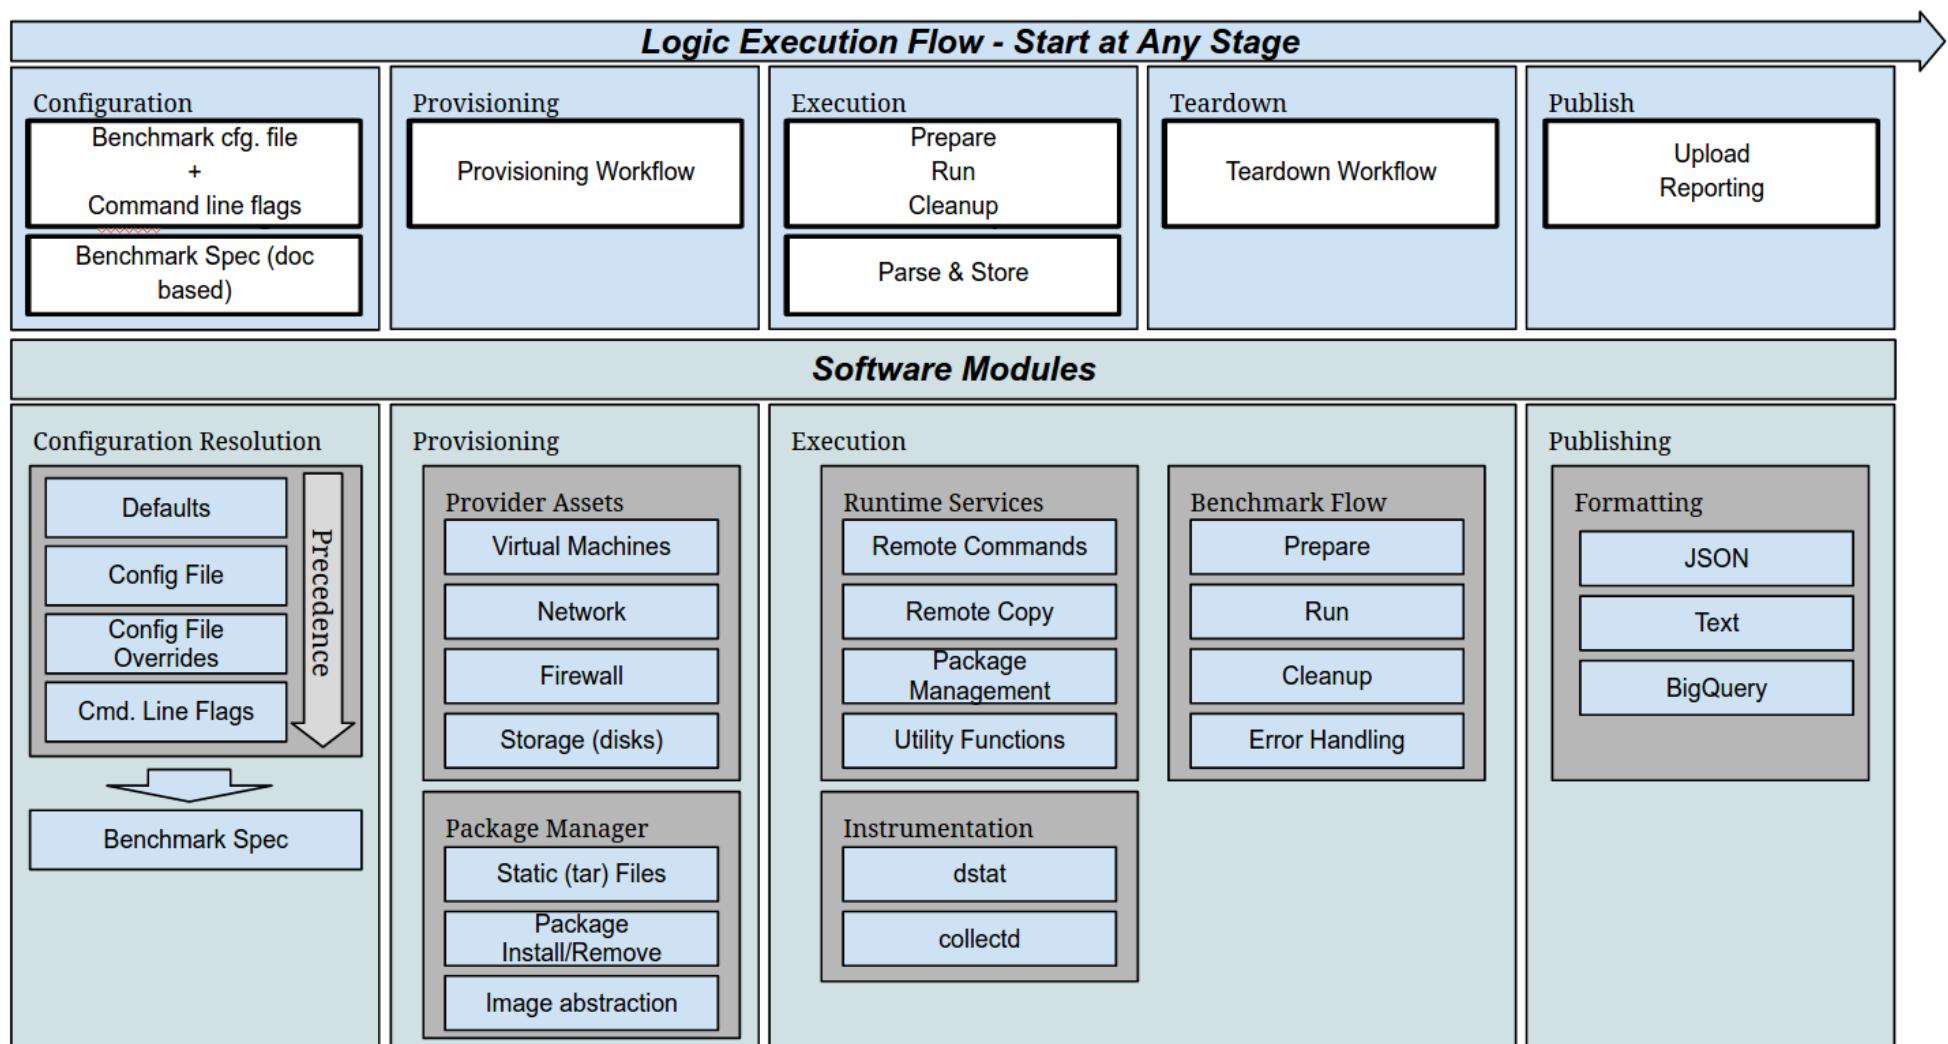
PerfKit Benchmarker

- Google tool
- Comparison of cloud solutions
 - Focus on network performance
- PKB features
 - Automatic setup and tear-down
 - Support for Docker and Open Stack
 - Most common workloads supported



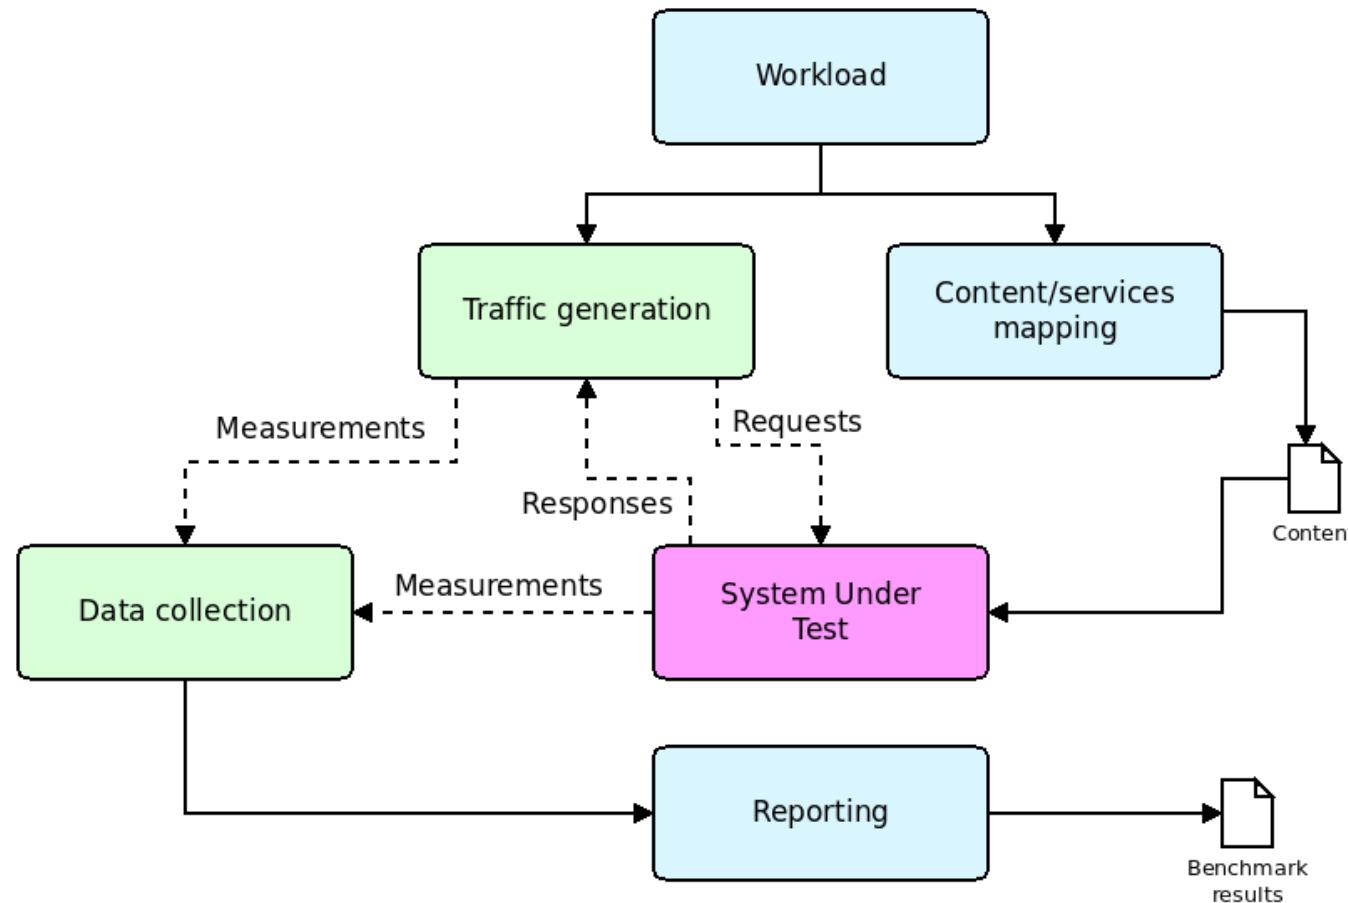
openstack
CLOUD SOFTWARE

PerfKit Benchmarker



Benchmark architecture

- Most examples seen share a common architecture
- Declined based on the goals of the performance study



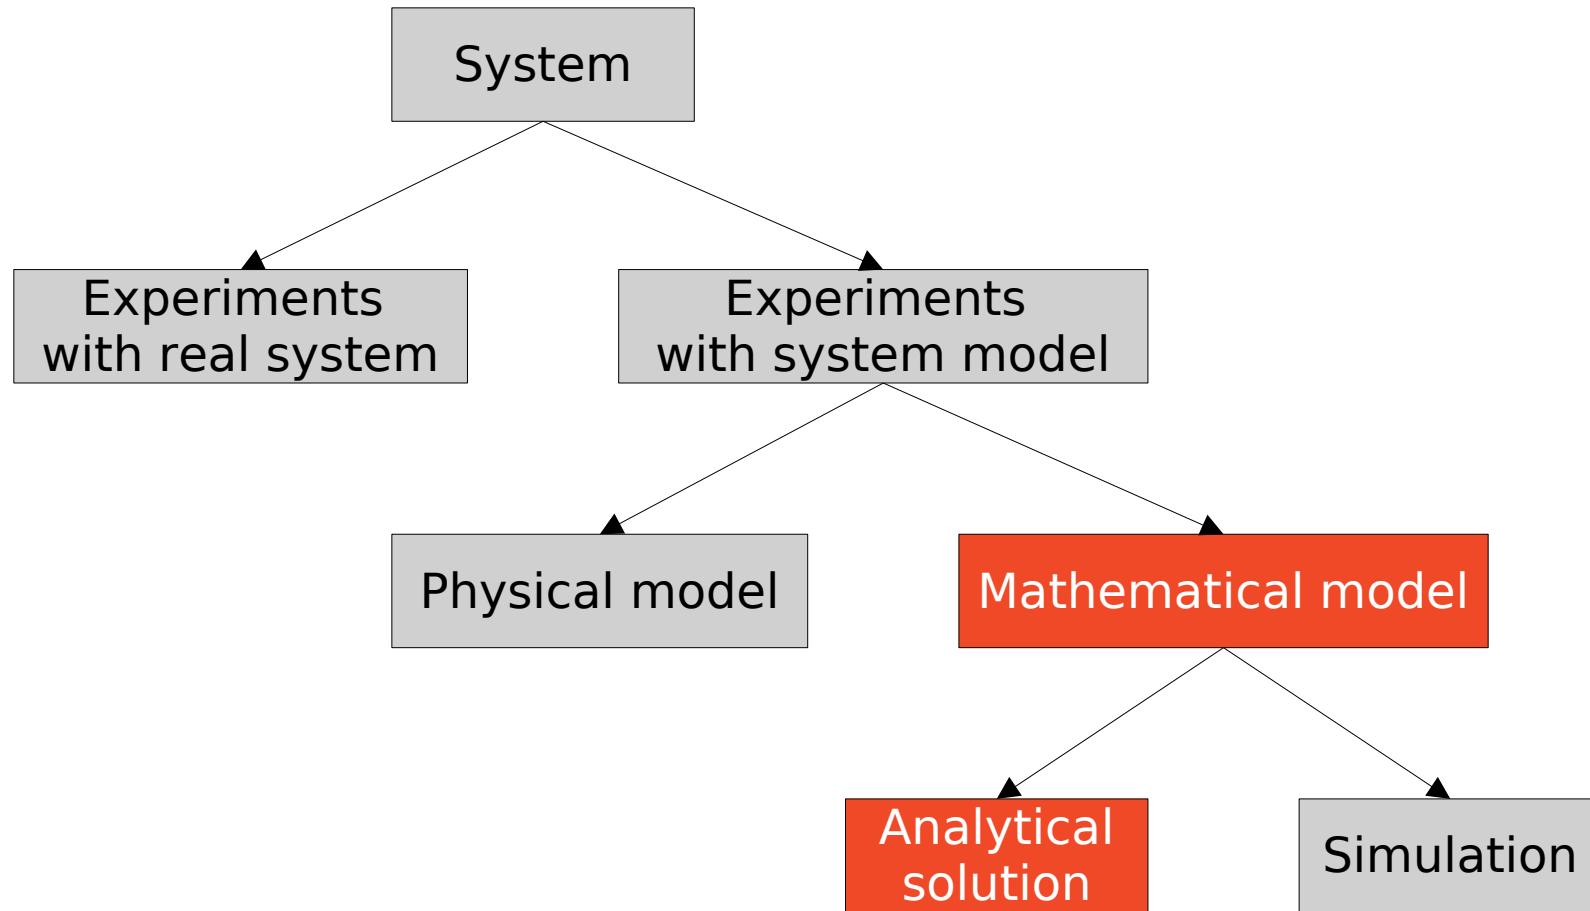
UNIMORE

UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA



Mathematical models

Possible approaches

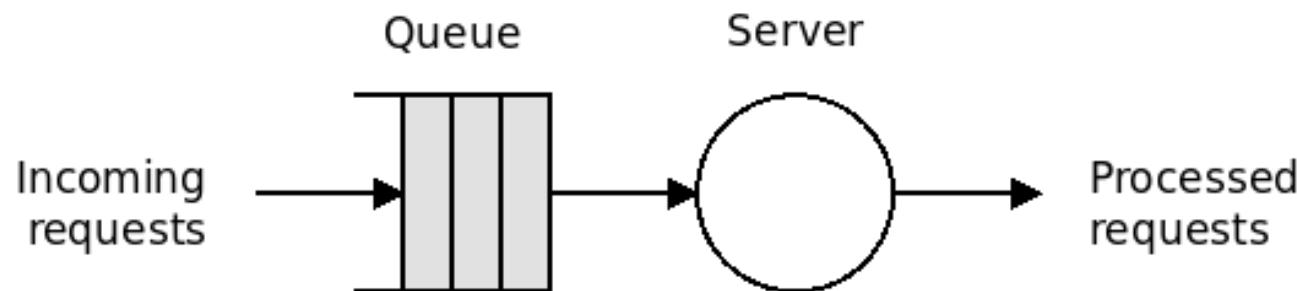


A simple example

- Multi-tier Web server
- Main contributions of response time:
 - Network
 - Front-end node
 - Middle-tier node
 - Back-end node
- Each node response time:
 - CPU time
 - Storage time (data retrieval)
 - Communication (LAN)

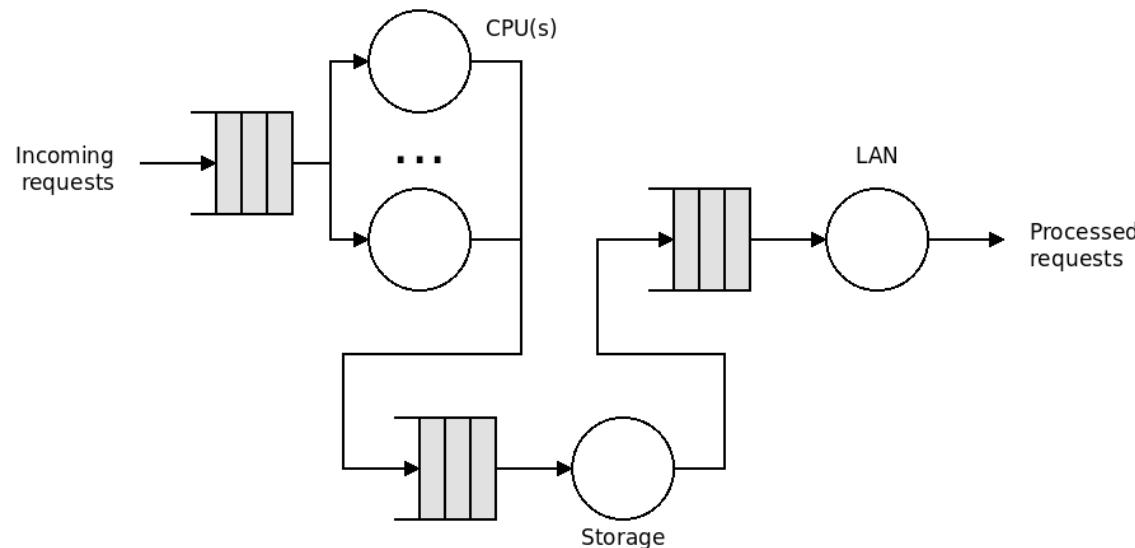
Network model

- Queuing model
- Example: **network**
 - Input: Incoming packets to be delivered
 - Output: delivered packets
 - Network is busy during transmission
 - Packets wait in routers queues



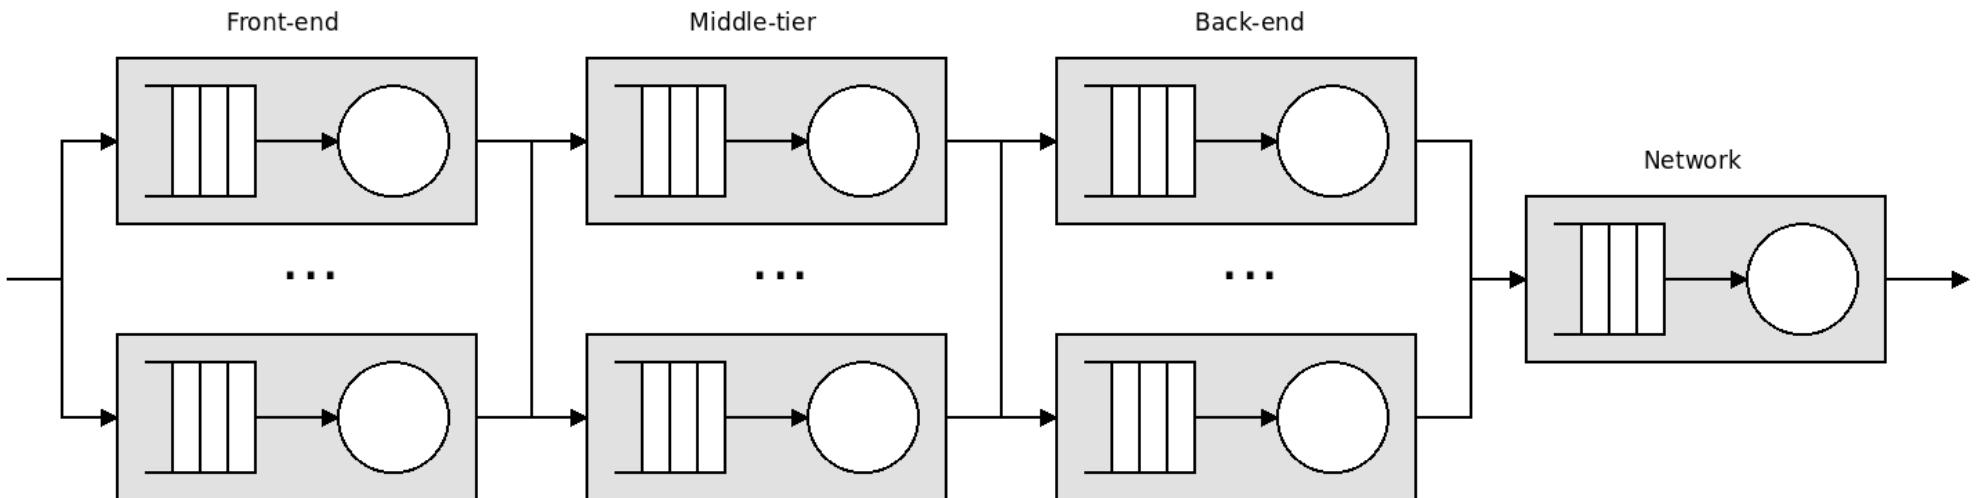
Nodes model

- Node model
 - CPU (multiple parallel servers for multi-core arch.)
 - LAN
 - Storage
- Each resource as a queue model



Big picture

- Complex model
- Several queues
- Same questions
 - Where is the bottleneck?
 - How to address it?



Modeling the problem

- Focus on a **single** node
 - Queues: CPU, storage, LAN
 - 1 server per each resource
- **Workload**
 - S_i : average service time at i-th resource ($1/\mu$)
 - W_i : average waiting time at i-th queue
 - λ_i : arrival rate at i-th queue

Modeling the problem

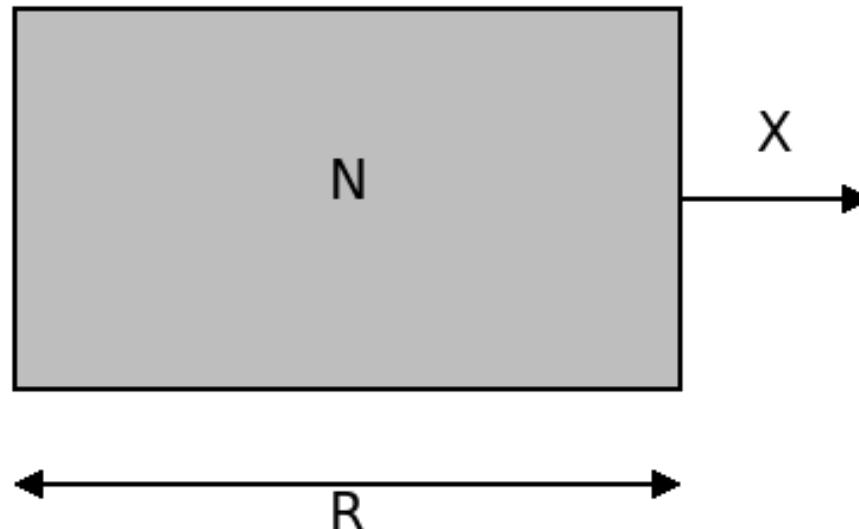
- **Performance metrics**
 - R_i : average response time at i-th queue
 - X_i : throughput of i-th queue
 - X_0 : overall system throughput
- **Number of jobs** at a resource
 - N_{iw} : requests in waiting queue at i-th resource
 - N_{is} : number of requests being serviced [0,1] (or [0,m])
 - N_i : number of requests at i-th resource

Utilization law

- B_i the time the server is busy in observation period τ
- C_i completed requests in period τ
- Utilization $U_i = B_i/\tau$
 - But $X_i = C_i/\tau$
 - Hence $1/\tau = X_i/C_i$
 - $U_i = B_i * X_i/C_i = (B_i/C_i) * X_i = S_i * X_i$
- In a steady state: $X_i = \lambda_i * S_i$
 - Often used notation: $\rho = \lambda/\mu$
- In a multiple server queue (with m servers)
 - $U_i = (X_i/m) * S_i$

Little's law

- System as a black box
 - N requests inside
 - X rate of outgoing requests (throughput)
 - R time spent within the system

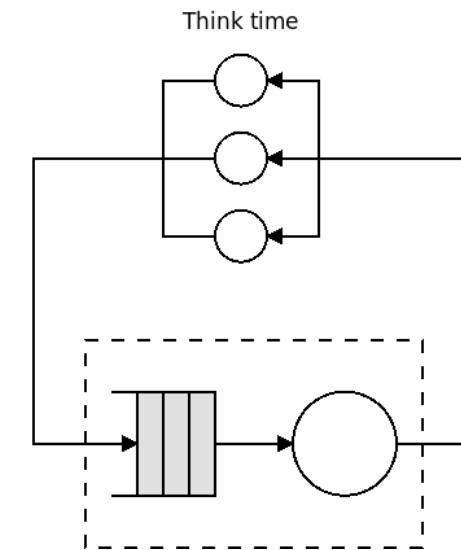
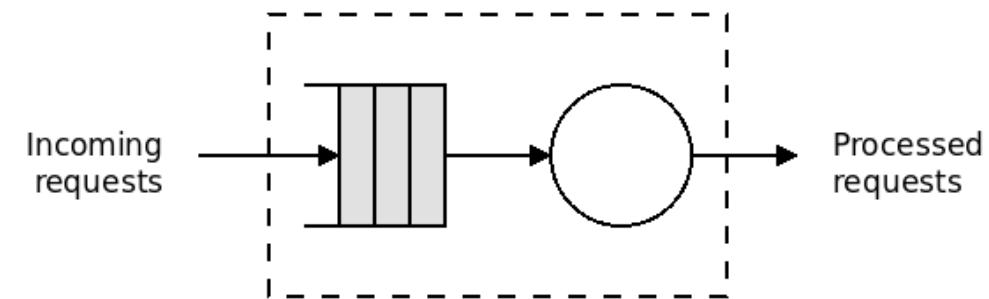


Little's Law

- General formulation:
 - $N = X * R$
- Applied to our problem:
 - Waiting: $N_{iw} = X_i * W_i$
 - Service: $N_{is} = X_i * S_i$
 - Response: $N_i = X_i * R_i$

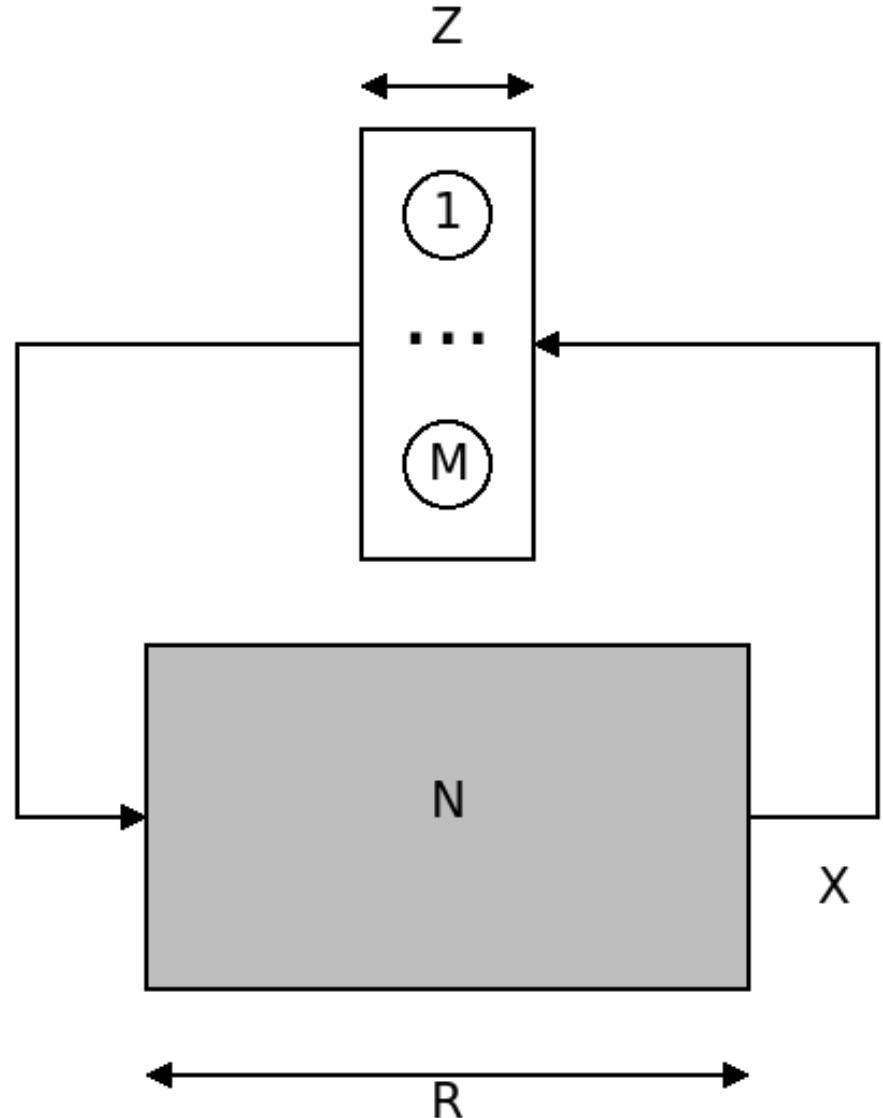
Open/Closed loop models

- Open loop:
 - Potentially infinite population of users
 - Requests come in and go out from system
- Closed loop:
 - Finite population of users
 - Requests circle through the system
 - Think time of users



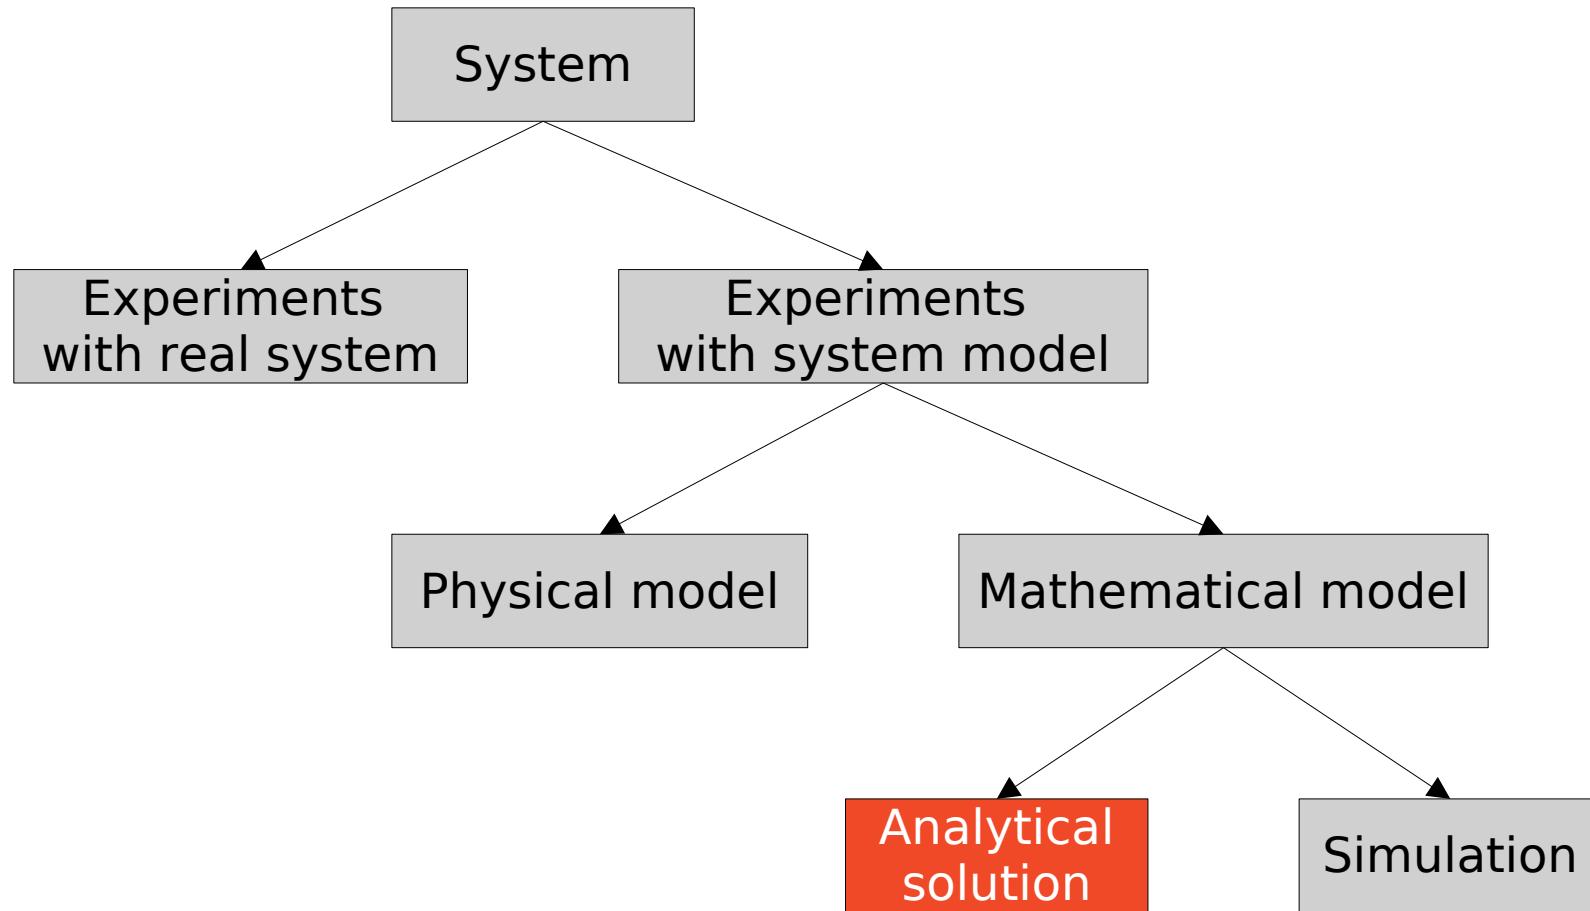
Response time law

- Overview:
 - System with a population of M users
 - Think time Z
- Closed loop formulation of little's law
 - $R+Z = M/X$



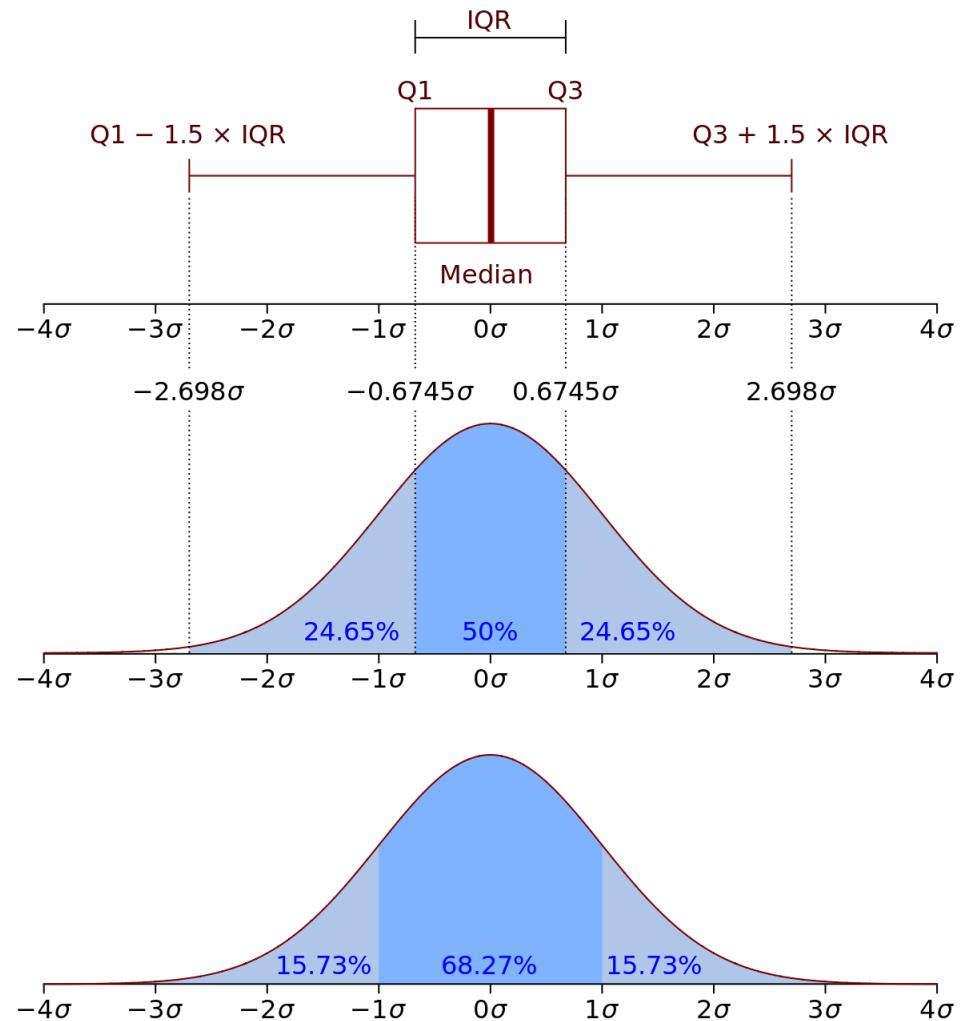
Queuing theory

Possible approaches



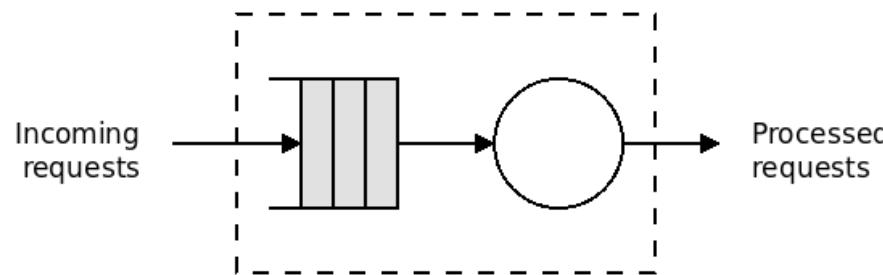
Stochastic process

- **Stochastic process**
 - A process driven by random behavior
 - Defines a family of random (aleatory) variables
- Aleatory variable
 - Can be described by its **statistical properties**
 - e.g., PDF



System overview

- Open loop
- Incoming job
 - Stochastic process
 - Inter-arrival time is a random variable
- Processed requests
 - Stochastic process
 - Processing time is another random variable

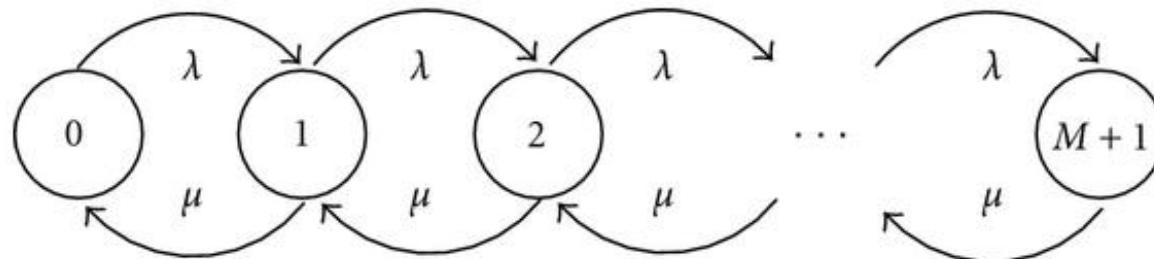


Some definitions...

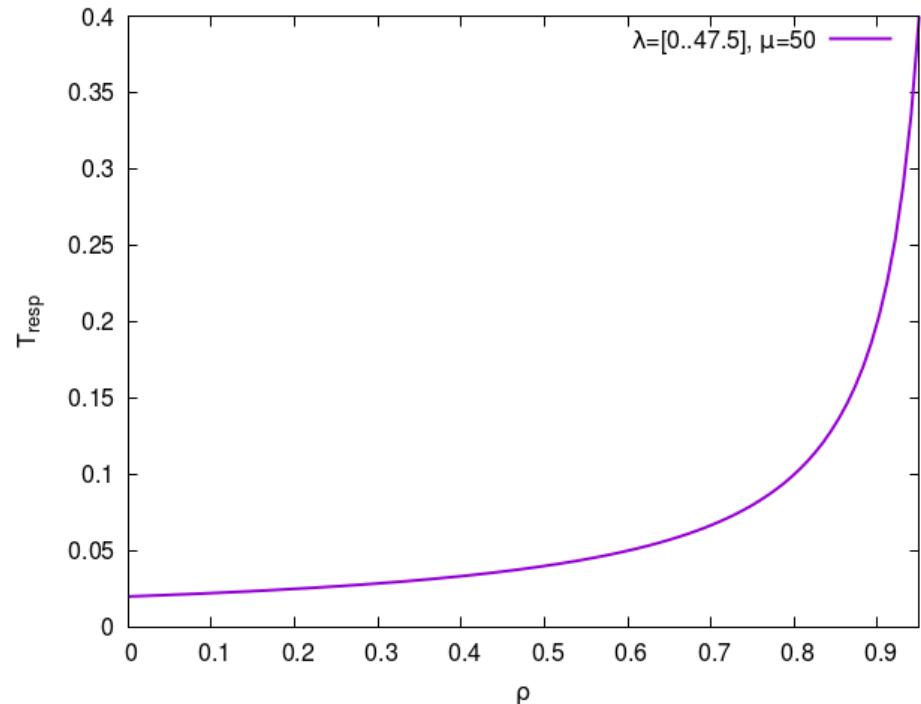
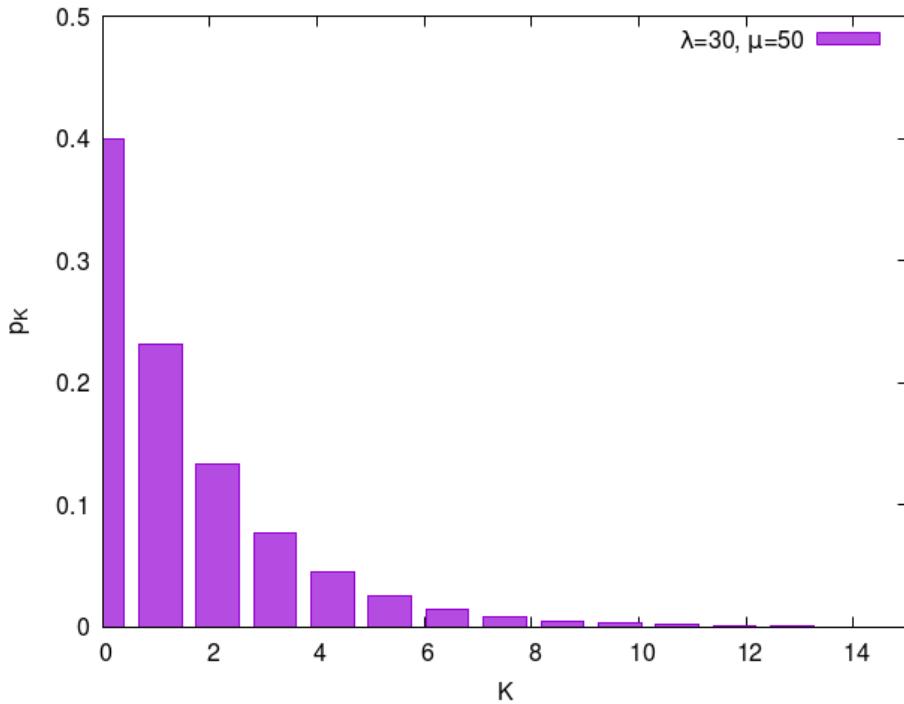
- System definition
 - {Arrival}/{Process}/{# Servers}/{Queue Len}...
 - e.g.: M/M/1/ ∞ (M/M/1 for short)
- Some stochastic processes:
 - M → Poisson (memory-less)
 - D → Deterministic
 - E_K → Erlang distribution
 - G → Generic
- Poisson process are easy to model
 - Use of **Markov Chain** (the M in the process)

Poisson process

- Status:
 - Elements in system
- Poisson process:
 - Rates do not depend on status
 - Arrival rate: λ
 - Leave rate: μ
 - Utilization: $\rho = \lambda / \mu$
 - Probability to be in state k : $p_k = (1 - \lambda / \mu) (\lambda / \mu)^k = (1 - \rho) \rho^k$



Poisson process



- Overview of Poisson process:
 - Infinite queue size
 - Probability of being state K
 - Response time as function of arrival rate: $T_{\text{resp}} = 1/(\mu - \lambda)$

Other useful formulas

- PASTA Theorem (M/G/1 systems)
- Pollaczek-Khinchin Formula

$$\begin{cases} \mathbb{E}[W] = \frac{\lambda \mathbb{E}[S^2]}{2(1 - \rho)} & = \frac{1 + C_v^2}{2} \cdot \frac{\rho}{1 - \rho} \cdot \mathbb{E}[S] \\ \mathbb{E}[T] = \mathbb{E}[S] + \frac{\lambda \mathbb{E}[S^2]}{2(1 - \rho)} & = \left(1 + \frac{1 + C_v^2}{2} \cdot \frac{\rho}{1 - \rho}\right) \cdot \mathbb{E}[S] \end{cases}$$

- $\mathbb{E}[W]$: average waiting time
- $\mathbb{E}[T]$: average response time
- $\mathbb{E}[S]$: average service time ($1/\mu$)

Other useful formulas

- Allen-Cuneen approximation (**G/G/N systems**)

$$W_M = \frac{P_{cb,N}}{\mu N(1 - \rho)} \left(\frac{C_S^2 + C_D^2}{2} \right), \quad P_{cb,N} \approx \begin{cases} (\rho^N + \rho)/2 & \text{if } \rho \geq 0.7 \\ \rho^{\frac{N+1}{2}} & \text{otherwise.} \end{cases}$$

- W_M Average waiting time
- $P_{cb,N}$ Probability that all servers are busy
- C_S, C_D , Coefficient of variation
 - arrival (C_S)
 - Service (C_D)
 - Coefficient of variation = Std. Dev / Average

A dirty trick (reprise)

- Let's turn Mr. Hyde for an example
- We want to sell a new CPU
 - Slightly faster than competitors
 - Expensive
- Need good advertisement
 - Pay 50% more for a 10% performance increase
- We can *tune* our benchmark
 - When system is heavily loaded small changes in power make the difference in response time
 - Pay 50% more for 2× faster



A dirty trick (reprise)

- We need to find the right load λ^*
 - Goal: **2 \times gain on response time**
- Real advantage
 - Throughout speedup is **K**
 - Competitor processor: 50 req/s
 - Our processor: 55 req/s $\rightarrow K=1.1$
- Assume a simple setup:
 - Poisson request arrival
 - Poisson service

Challenge

- How to model the problem?
- How to find λ^* ?
- **Volunteers?**

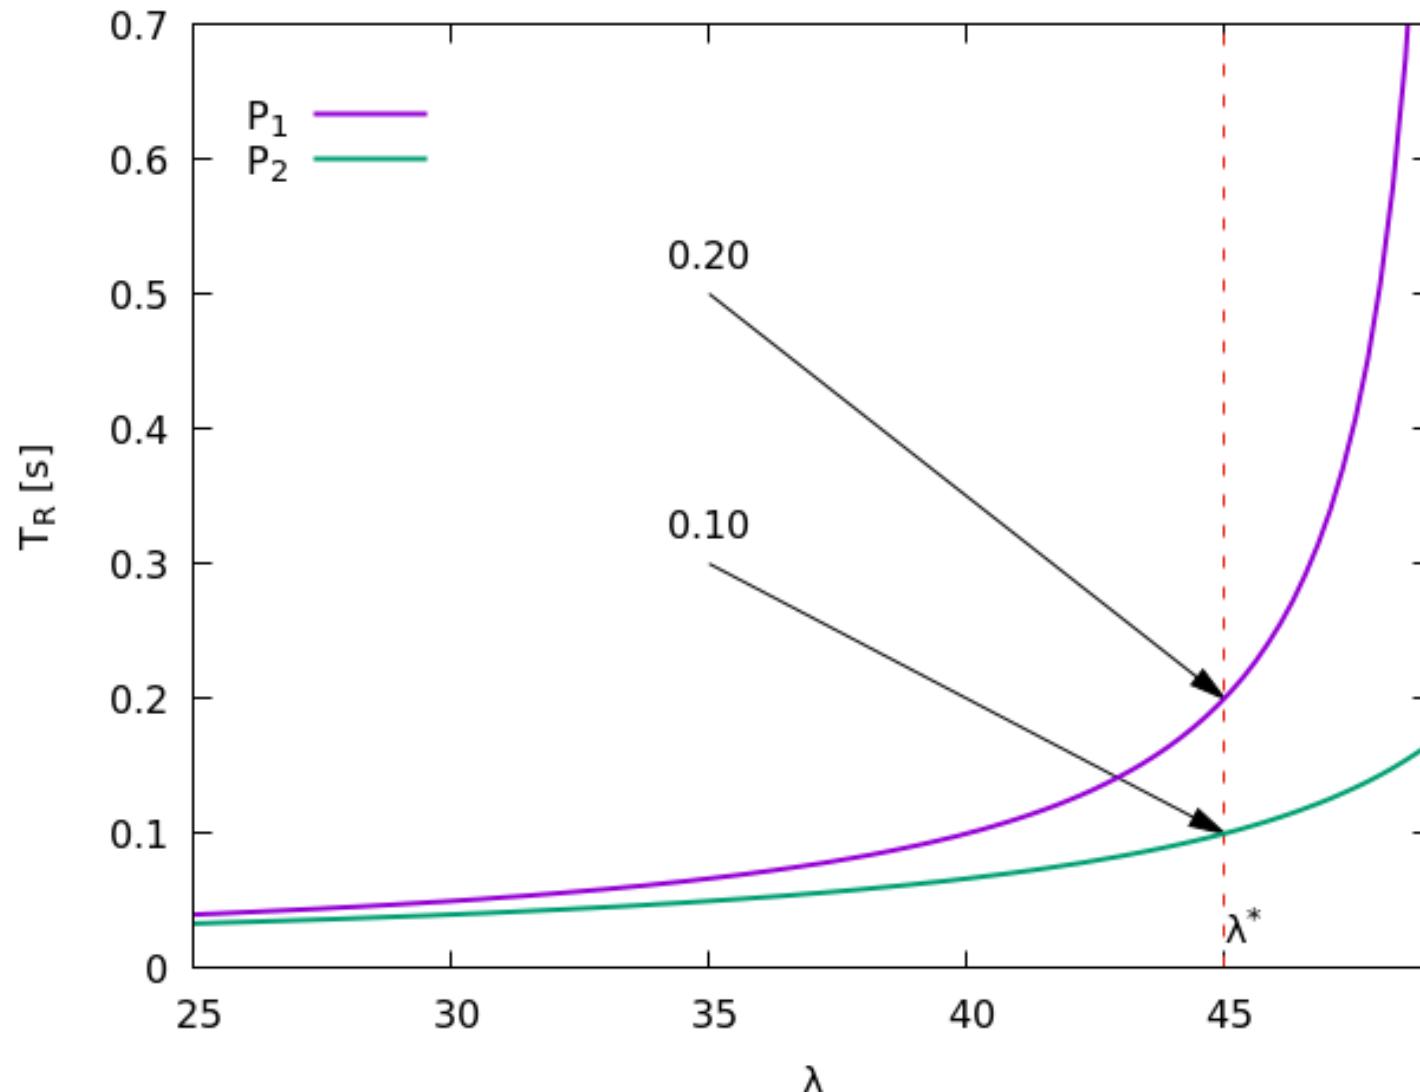
A dirty trick (reprise)

- We model the systems as M/M/1 systems
- Processor P1 (slow):
 - $\mu_1 = \mu$
 - $T_1(\lambda) = 1/(\mu - \lambda)$
- Processor P2 (our, fast):
 - $\mu_2 = K \mu$
 - $T_2(\lambda) = 1/(K\mu - \lambda)$
- P2 is 2x faster at λ^*
 - $T_1(\lambda^*) = 2 T_2(\lambda^*)$

A dirty trick (reprise)

- Solution is easy:
 - $\lambda^* = (2-K)\mu$
- We can now plot the results

A dirty trick



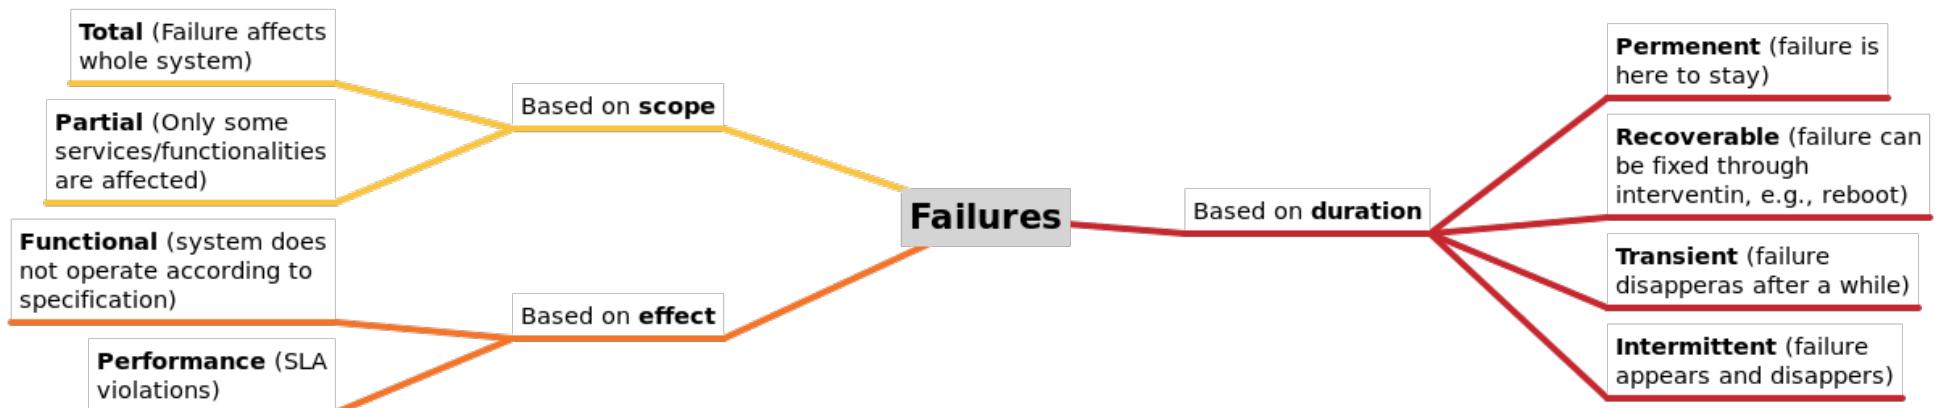
Other exercises

- We have an incoming load $\lambda > \mu$
- We want to scale-up our system → K servers
- Assumption:
 - Requests can be distributed uniformly over servers
- How many servers do I need to avoid overload?
- Assumption
 - I have a maximum response time T_{SLA}
- How many servers do I need to meet the SLA?
- Volunteers welcome

Failures

Availability

- Being available → not being in a failure state
- Classification of failures
 - Several criteria



Failures

- Classification based on Duration
- Permanent
 - Failure is here to stay
- Recoverable
 - It is possible to return online after intervention (e.g., restart, reboot)
- Transient
 - Failure disappears after a while
- Intermittent
 - Failure that appears and disappears

Failures

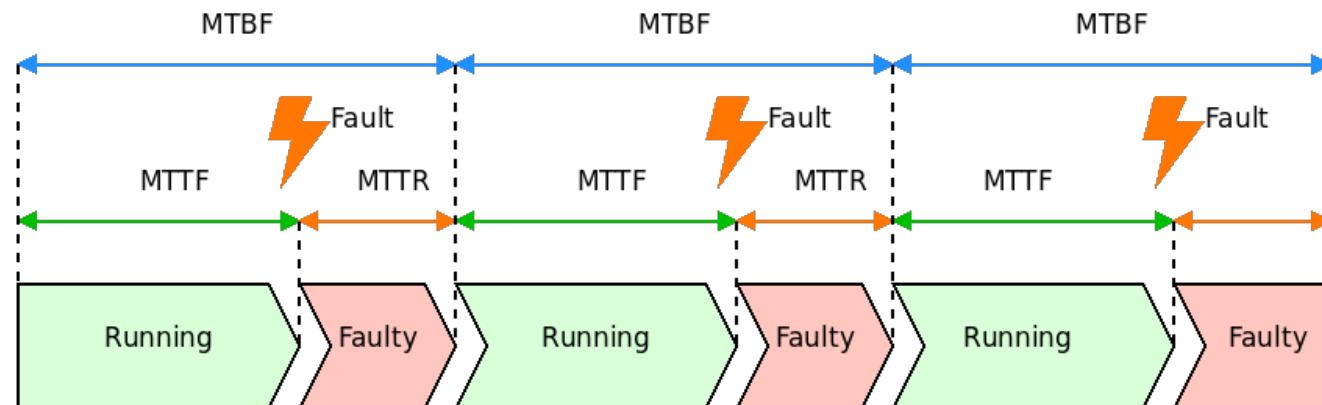
- Classification based on Effect
- **Functional**
 - The system does not operate according to its functional specifications
 - “System is not working right”
- **Performance**
 - System works but not in a timely fashion
 - SLA violations

Failures

- Classification based on Scope
- Total
 - Failure affects the whole system
- Partial
 - Only some services/functionalities are affected

Failure metrics

- **MTBF**
 - Mean Time Between Faults
- **MTTF**
 - Mean Time To Fault
- **MTTR**
 - Mean Time To Repair
- **MTBF=MTTF+MTTR**



Availability classes

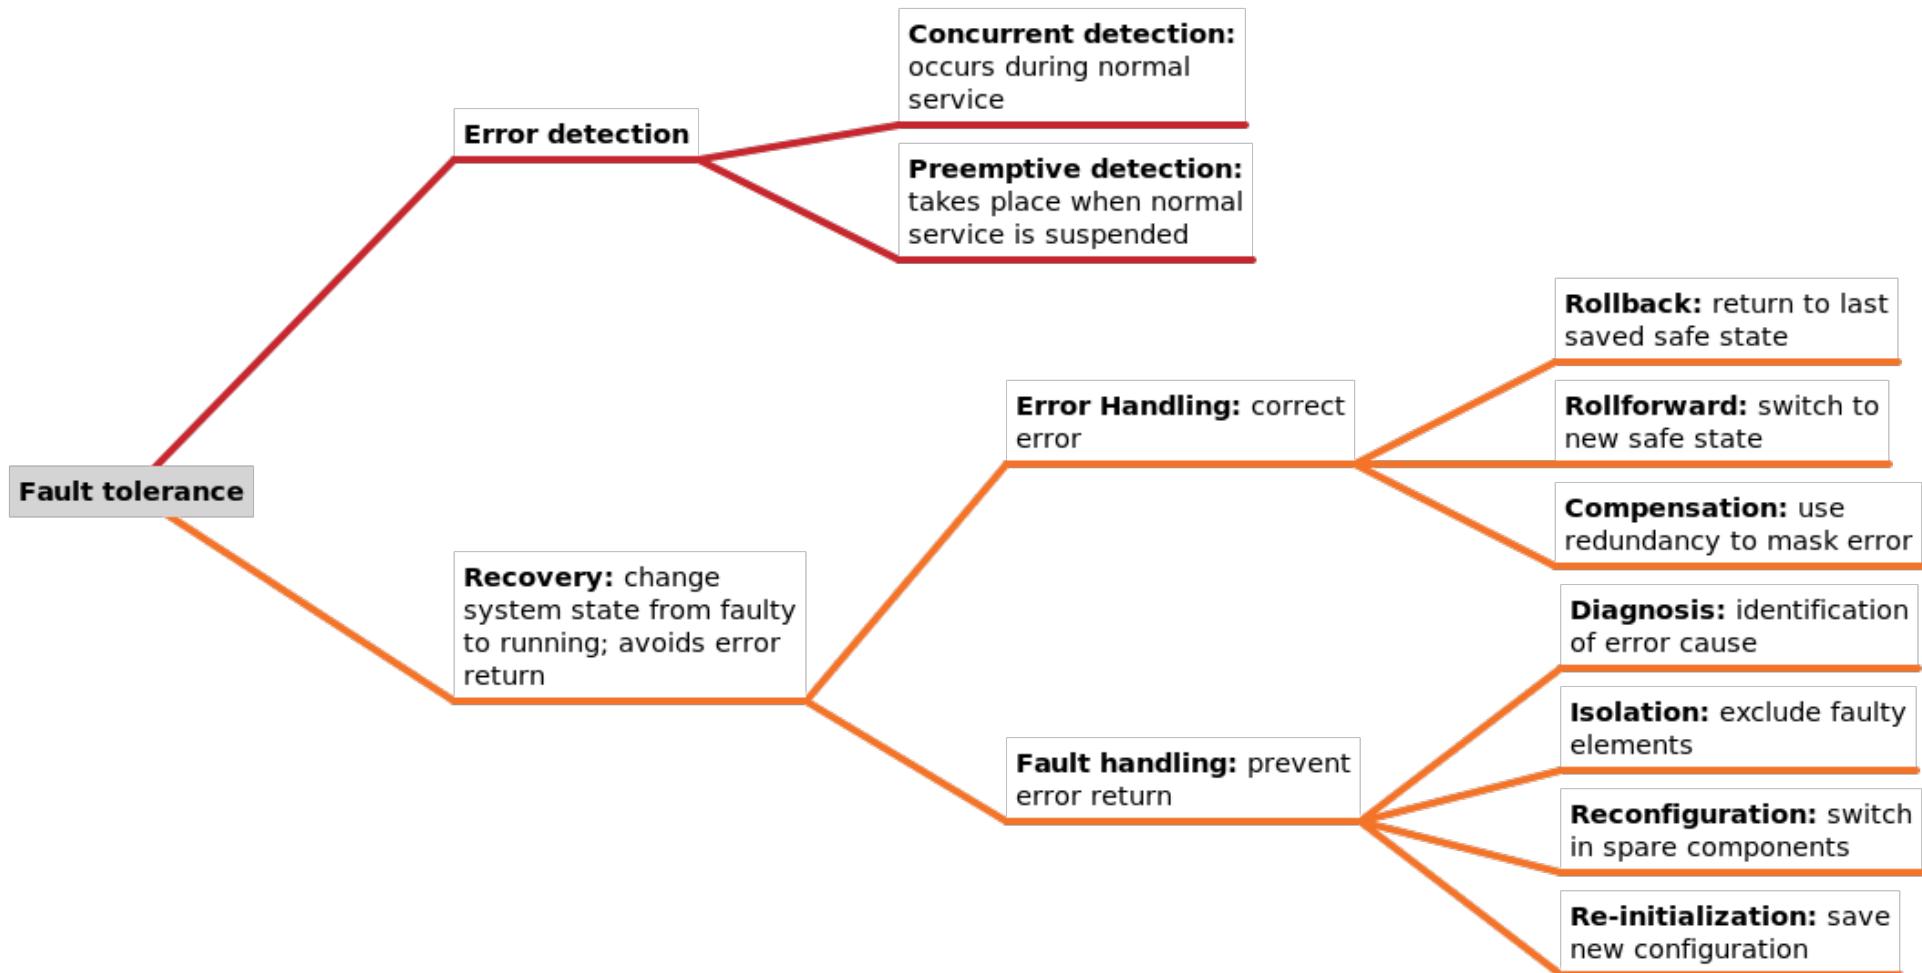
Availability Class	Availability	Unavailable [min/y]	System Type
1	90.0%	52560 (36d)	Unmanaged
2	99.0%	5256 (3.6d)	Managed
3	99.9%	526 (8.8h)	Well-managed
4	99.99%	52.6	Fault-tolerant
5	99.999%	5.3	Highly Available
6	99.9999%	0.53	Very Highly Available
7	99.99999%	0.053	Ultra Available

- **Availability**
 - Fraction of time when the system is operational
 - Availability $A=MTTF/MTBF$

Availability and reliability

- Availability
 - Probability that a system is available at time t
- Reliability
 - Probability that a system works correctly during a period τ
- If the observation period τ is long enough **Availability and Reliability converge** to the same value

Fault-tolerance checklist



Some figures

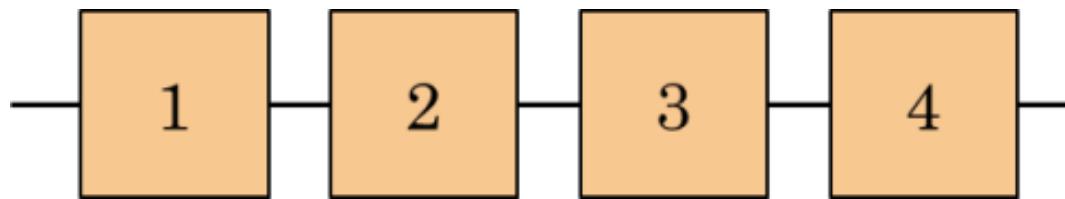
- Data Center problems in one year
 - ~0.5 overheating (power down most machines in <5 mins, ~1-2 days to recover)
 - ~1 PDU failure (~500-1000 machines suddenly disappear, ~6 hrs to come back)
 - ~1 rack-move (plenty of warning, ~500-1000 machines powered down, ~6 hrs)
 - ~1 network rewiring (rolling ~5% of machines down over 2-day span)
 - ~20 rack failures (40-80 machines instantly disappear, 1-6 hours to get back)
 - ~5 racks go wonky (40-80 machines see 50% packet loss)
 - ~8 network maintenances (might cause ~30-minute random connectivity losses)
 - ~12 router reloads (takes out DNS and external vIPs for a couple minutes)
 - ~3 router failures (have to immediately pull traffic for an hour)
 - ~dozens of minor 30-second blips for dns
 - ~1000 individual machine failures (2-4% failure rate, machines crash at least twice)
 - ~thousands of hard drive failures (1-5% of all disks will die)

Combining reliability

- Given N sub-systems
 - Each subsystem I has a failure probability r_i
- What is the overall failure probability?
 - Depends on how systems interacts
- **Serial** systems
- **Parallel** systems

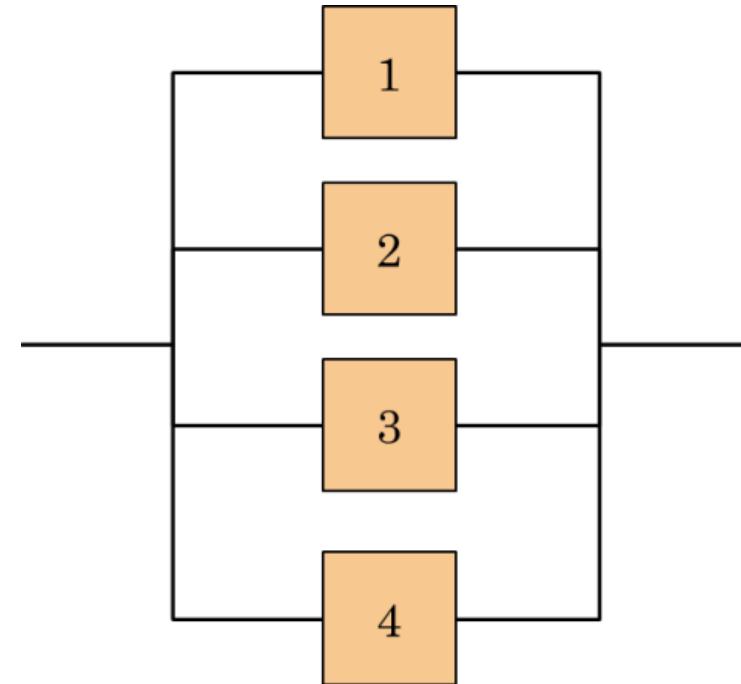
Serial systems

- Global failure → one sub-system failed
 - Like Christmas lights
- System reliability:
 - $R = \prod r_i$
- System fails if one sub-system fails
- System is available if every sub-system is available



Parallel systems

- Global failure → every sub-system failed
- System reliability:
 - $R=1-\prod(1-r_i)$
- System is available if at least one sub-system is available



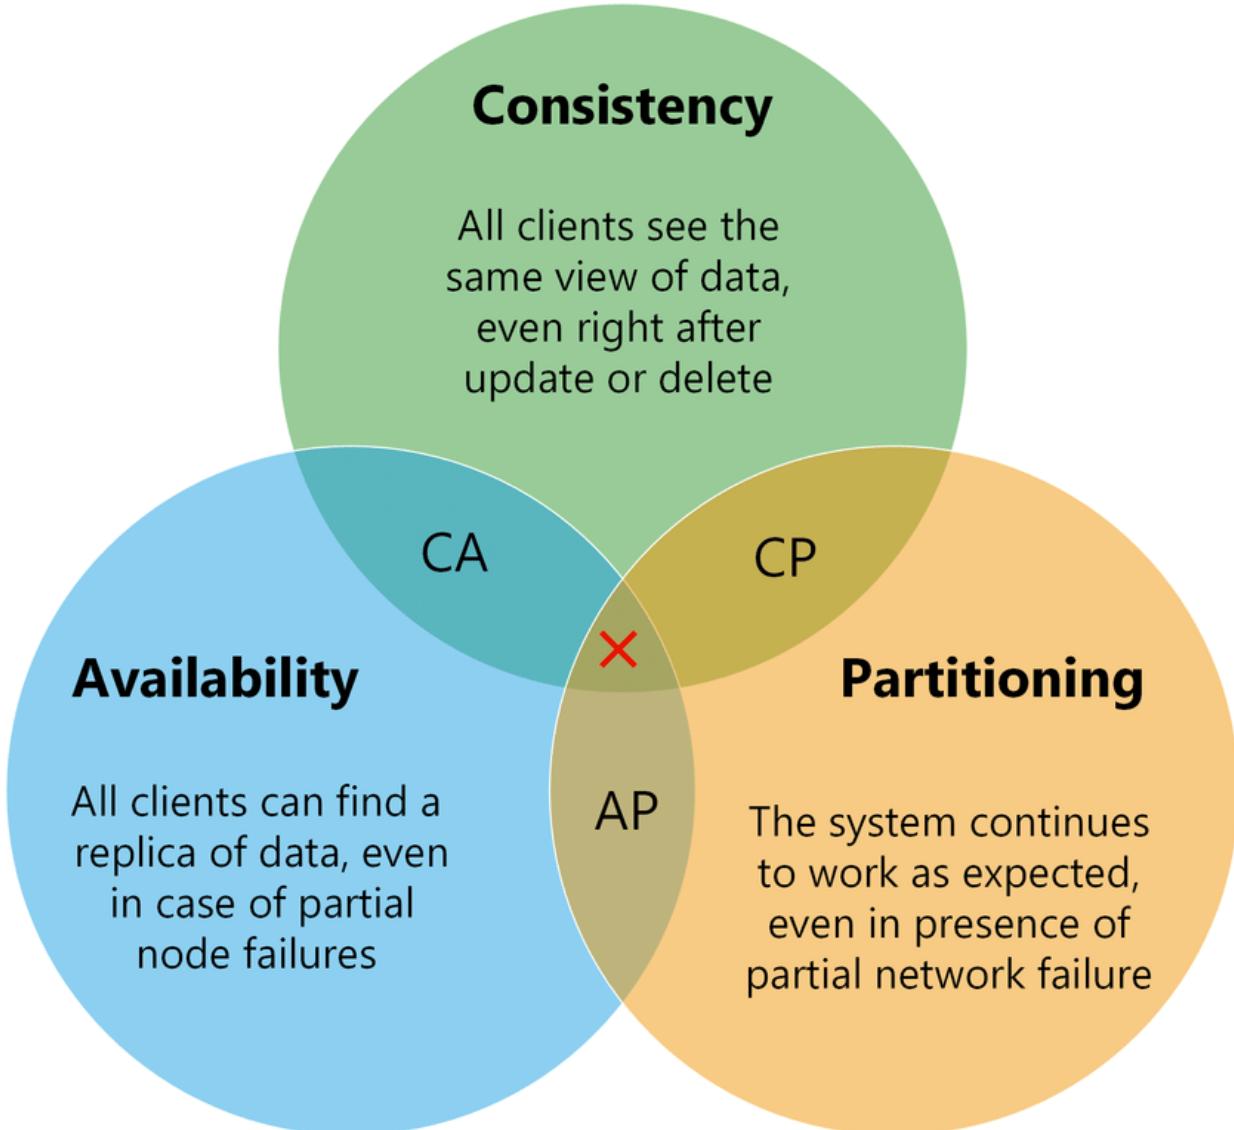
Some questions

- Given subsystem with identical reliability
- In a parallel system
 - How many subsystem do I need to guarantee a reliability of X?
- In a serial system
 - What is the maximum number of subsystem to guarantee a reliability of X?
- **Volunteers?**

CAP Theorem

- **CAP**
 - Consistency/Availability/Partition-tolerance
- **Consistency:** Every read return data from most recent write
- **Availability:** Every request returns a non-error response
- **Partition-tolerance:** The system works also in case of network partitions (messages lost or delayed)
- CAP theorem
 - Can choose only **2 out of 3** properties

CAP Theorem



Some useful hints

- Look out for **bottlenecks**
- Look out for **single points of failure**
- Keep it simple and stupid (**KISS**)
 - Centralized control is simple!
- If it is not **tested**, don't rely on it
- Testing is not easy
 - Hundreds of services
 - Loosely couples services
 - Thousand of machines
 - How to deal with it?



Additional information

- Some useful readings:
- D. Menascé, V. Almeida,
“Capacity Planning for Web Services”
- K. Trivedi, A. Bobbio,
“Reliability and Availability Engineering”
- K. Trivedi “Probability and Statistics with Reliability, Queuing, and Computer Science Applications”

Workload forecasting

Forecasting

- **Forecast:**
a statement of what is judged likely to happen in the future, especially in connection with a particular situation,
- Trying to understand how workload is **evolving**
- A critical part of capacity planning

A good forecast

- Not just a number but a **set of scenarios**
 - optimistic/average/pessimistic
- May involve **many parameters** of workload
 - Not just a change in workload intensity
- Examples:
 - Session arrival rate
 - Requests per session
 - Think time between interactions
 - Probability ratio in multiple services (workload mix)
 - Response size

Steps for forecasts

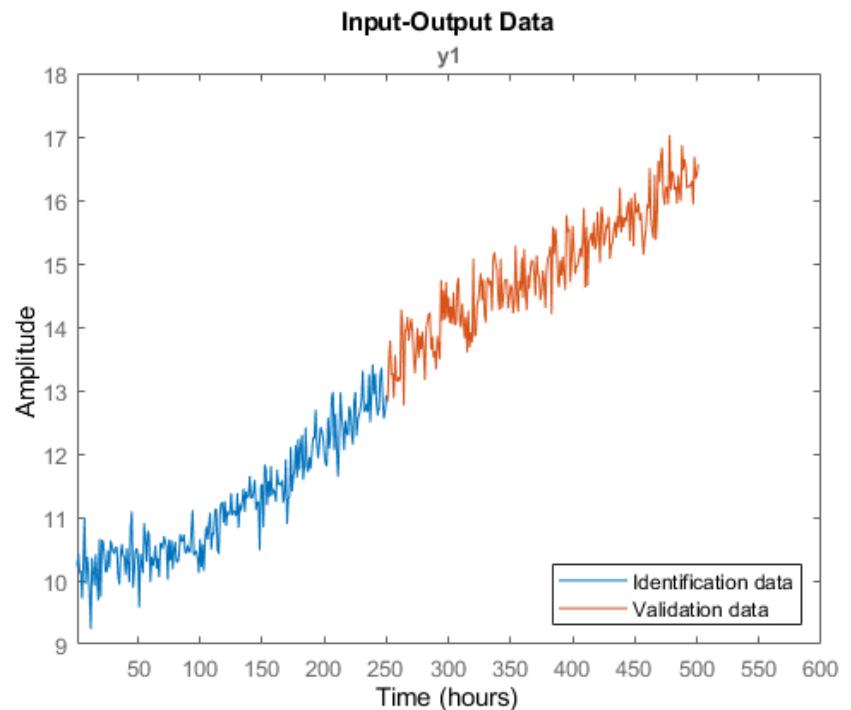
- Defining objectives (what are we trying to predict?)
- Select information sources
 - Interviews
 - Reports
 - Historical data
- Gather data (quantitative/qualitative)
- Build models
- Forecast
- Validation

Qualitative forecast

- Typically based on human sources
- Can be useful to build a reference scenario
- Examples:
 - Interviews with end users
 - Questionnaires to a large set of people
 - Collecting opinions of experts
 - More free-format interview
 - Analysis of technological trends
 - Access to specialized reports

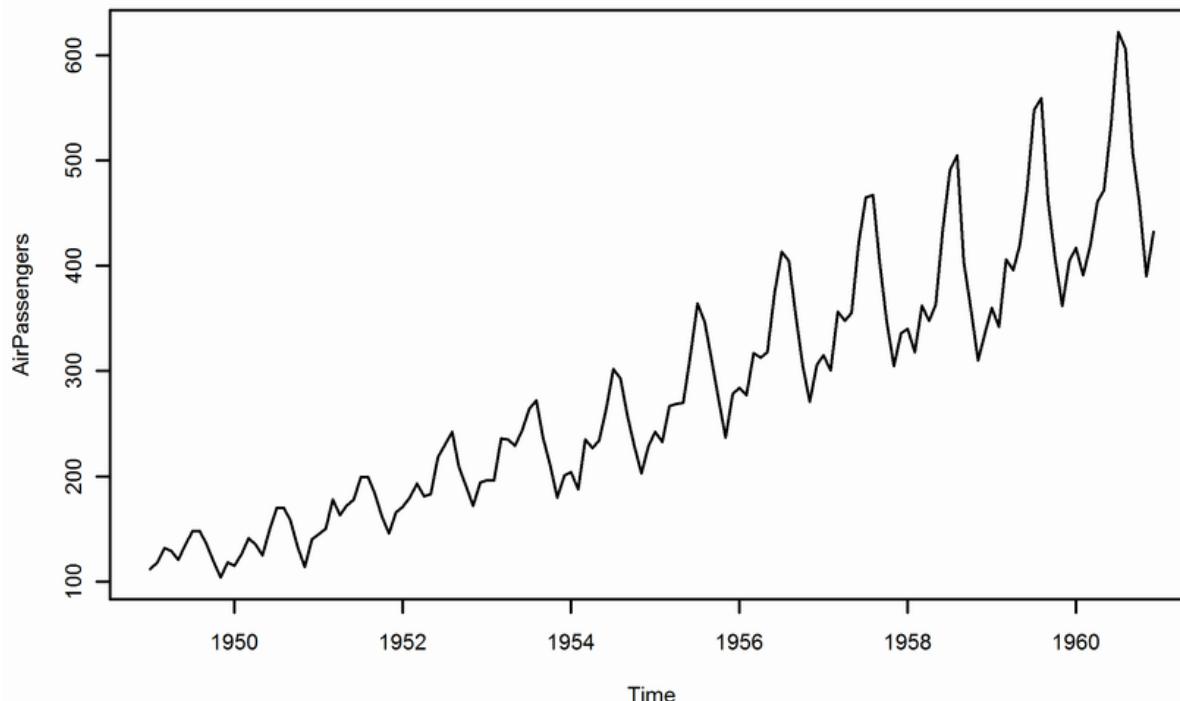
Validation

- Can use part of **historical data**
- Apply forecast to past
- It should describe the present
- Similar to Training/Validation process in ML



Historical data analysis

- Identify trends in time series
 - Linear trend
 - Periodic components



Error measures

- Mean error
 - $ME = \sum_t (Y_t - F_t) / n$
- Mean square error
 - $MSE = \sum_t (Y_t - F_t)^2 / n$
- Sum of square error
 - $SSE = \sum_t (Y_t - F_t)^2$

Linear regression methods

- **Interpolation function**
 - $Y = a + bX$
- **Minimize square error (Least Square method)**
 - $b = (n \sum_t (Y_t X_t) - (\sum_t X_t)(\sum_t Y_t)) / (n \sum_t (X_t)^2 - (\sum_t X_t)^2)$
 - $a = \text{avg}(Y) - b(\text{avg}(X))$

How predictable is a series?

- Typical test:
- Autocorrelation
 - Correlation between:
 - Time series $X(t)$ and
 - Time series $X(t-lag)$
 - Evaluation as a function of lag
- Slow autocorrelation decay:
 - easy to predict
- Fast autocorrelation decay:
 - highly irregular time series
- Spikes for a specific lag value
 - Periodic behavior

